

# CS 4342: Class 3

Jacob Whitehill

# Combining multiple predictors

- Determining smile/non-smile based on a single comparison is very weak.
- What if we combined *multiple* pairs and took the majority-vote (choose non-smile if tied) across all  $m$  comparisons?

$$g^{(j)}(\mathbf{x}) = \mathbb{I}[\mathbf{x}_{r_1, c_1} > \mathbf{x}_{r_2, c_2}]$$
$$\hat{y} = g(\mathbf{x}) = \mathbb{I} \left[ \left( \frac{1}{m} \sum_{j=1}^m g^{(j)}(\mathbf{x}) \right) > 0.5 \right]$$

# Step-wise classification

- Pseudocode:

```
predictors = [] # Empty list
```

```
For j = 1, ..., m:
```

1. Find next best predictor given what's already in predictors
2. Add it to predictors

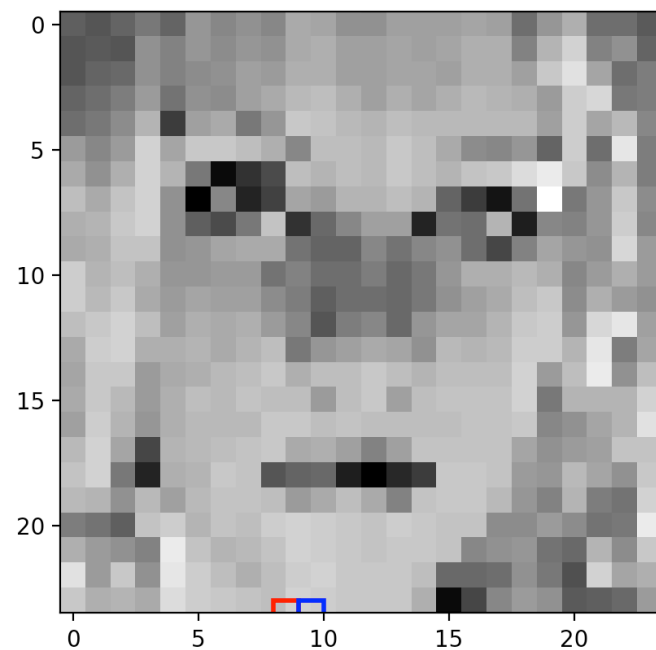
- Run smile\_demo.py and optimize on 10 images.

# Step-wise classification

- Accuracy (on 10 images): 100%.
- Learned feature:

# Step-wise classification

- Accuracy (on 10 images): 100%.
- Learned feature (somewhat counterintuitive):



- What happened?

# Overfitting

- When we optimized the  $m=1$  features on a set of just 10 images, we discovered a **spurious relationship** between the image  $\mathbf{x}$  and the target label  $y$ .
- Spurious: the relationship would not **generalize** to a much larger set of images.

# Overfitting

- When we optimized the  $m=1$  features on a set of just 10 images, we discovered a **spurious relationship** between the image  $\mathbf{x}$  and the target label  $y$ .
  - Spurious: the relationship would not **generalize** to a much larger set of images.
- **Problem:** we have many features (331200) but very few images (10) we need to classify.
  - Out of 331200, it's not hard to find a few features that happen to discriminate smiles/non-smiles just by chance.
- This is called **overfitting** to the dataset.

# Training versus testing data



# Training versus testing

- In machine learning, we always optimize the parameters/features of our classifier/regressor on a **training set**  $\mathcal{D}^{\text{train}}$ .
- We then measure accuracy on a **testing set**  $\mathcal{D}^{\text{test}}$  that is disjoint from (contains no common elements with) the training set.
- The accuracy on the testing set characterizes how well our machine will perform on *new* data.
- The training and testing sets should be collected in the *same manner*.

# Training versus testing

- What if the test accuracy was bad?
  - Then we should make some changes to the architecture or training procedures of our machine and re-train.
- To estimate the accuracy of the new machine, we should evaluate it on a *new* test set!
  - Why?

# Training versus testing

- When we re-train a ML model, there are many things we can change, e.g.:
  - L1/L2 regularization strength
  - SVM kernel type + parameters
  - Number of neural network layers, #units/layer
  - Learning rate, momentum, etc.

# Training versus testing

- When we *iteratively* re-train many ML models by optimizing on the *test set*, we might find good values for these choices just *by chance*.

# Implicit cheating

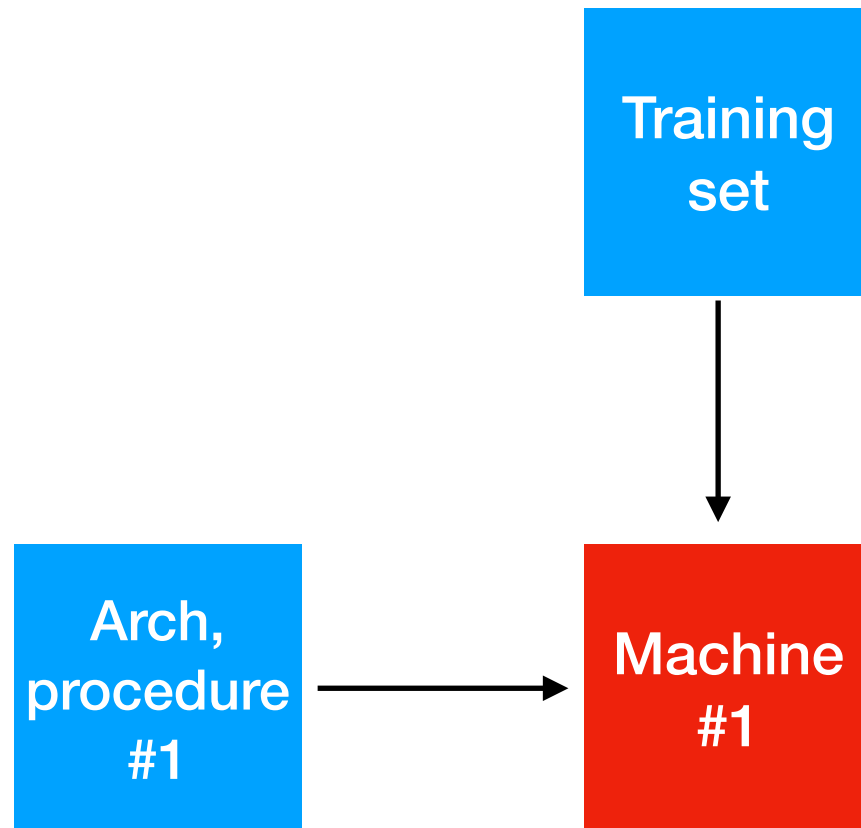
- Train a model on the training set using a particular architecture and training procedure.

Training  
set

Arch,  
procedure  
#1

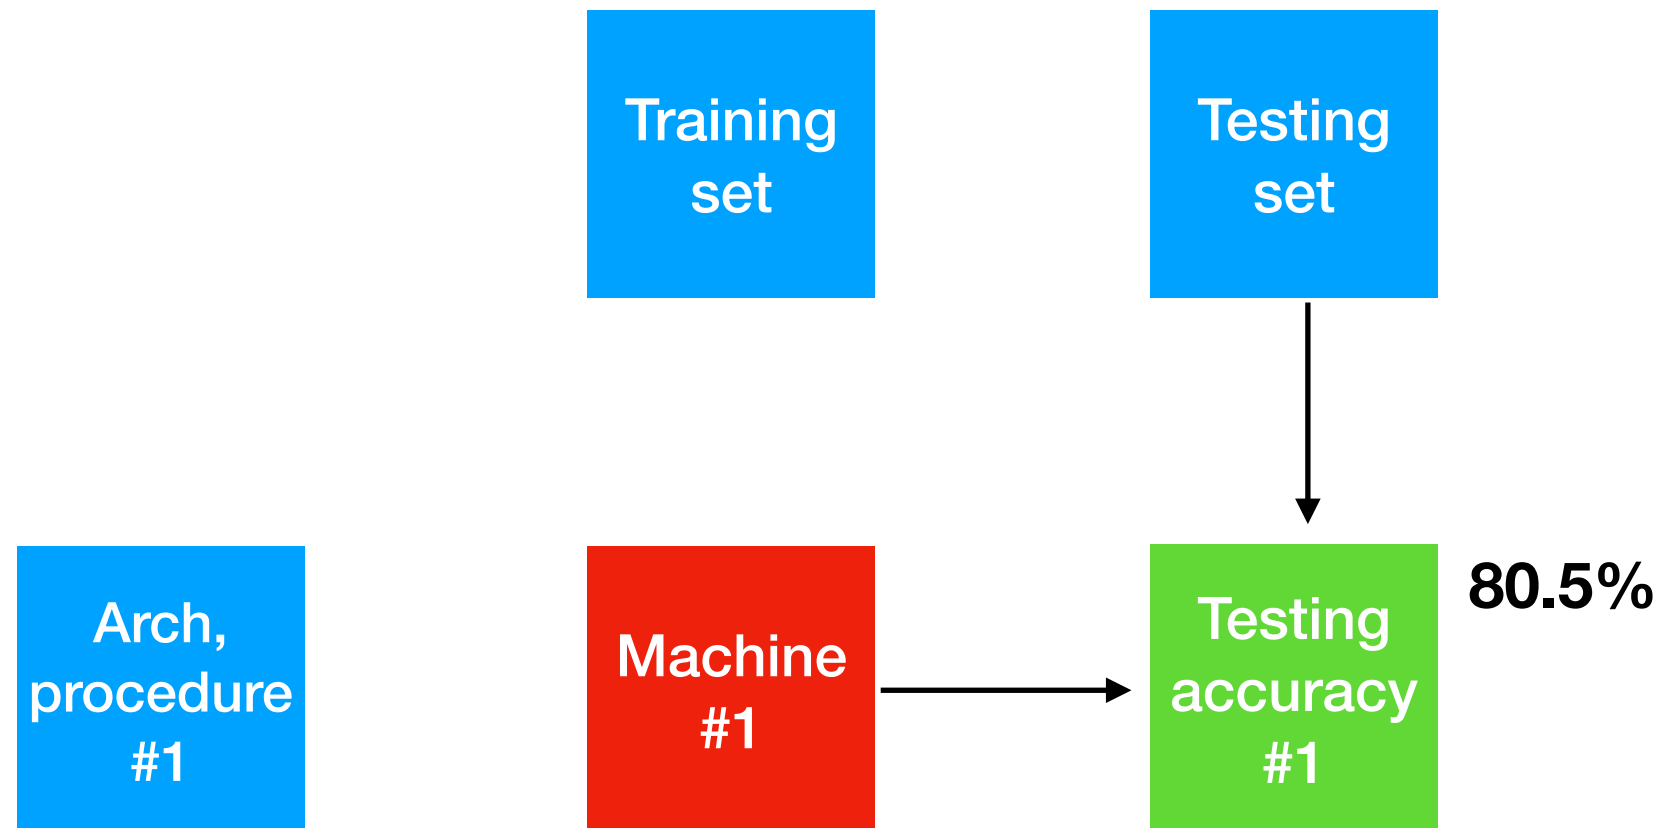
# Implicit cheating

- Train a model on the training set using a particular architecture and training procedure.



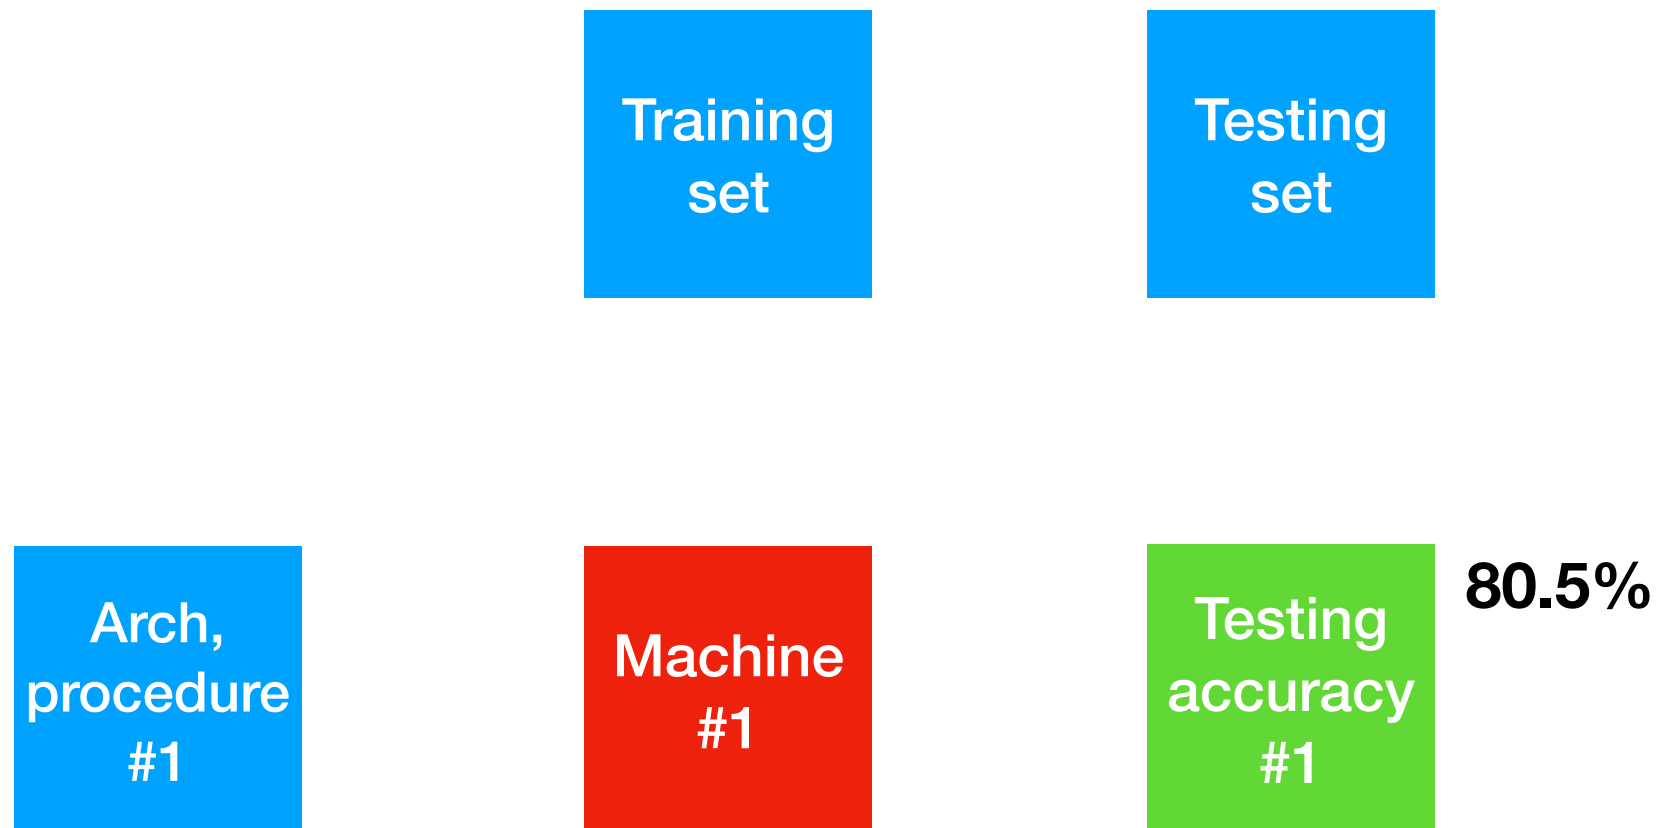
# Implicit cheating

- Evaluate the trained machine on the testing set.



# Implicit cheating

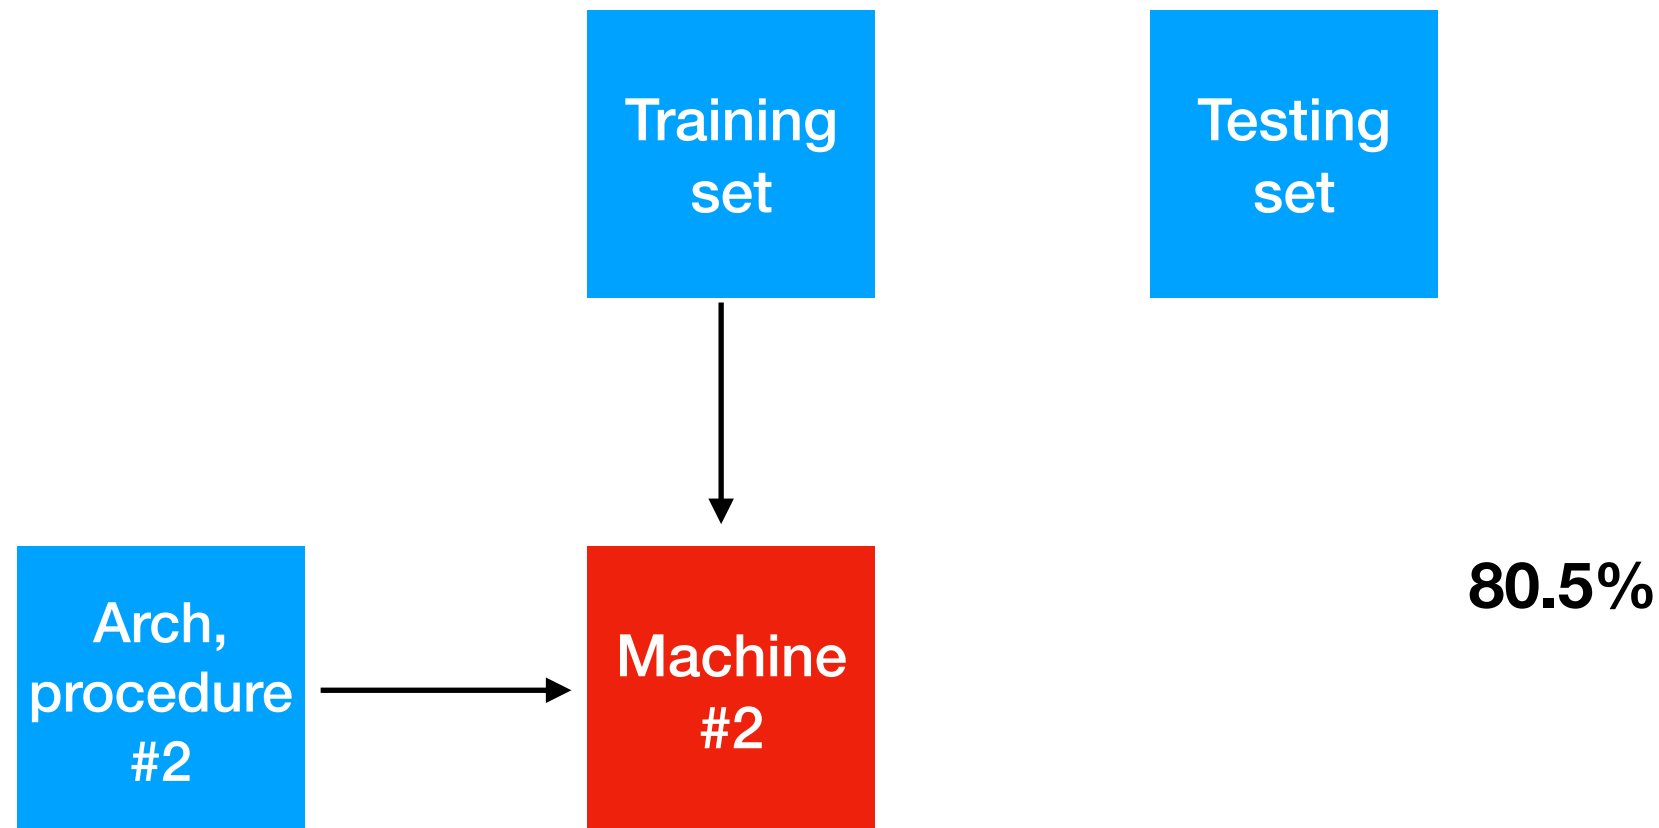
- Accuracy not good enough?





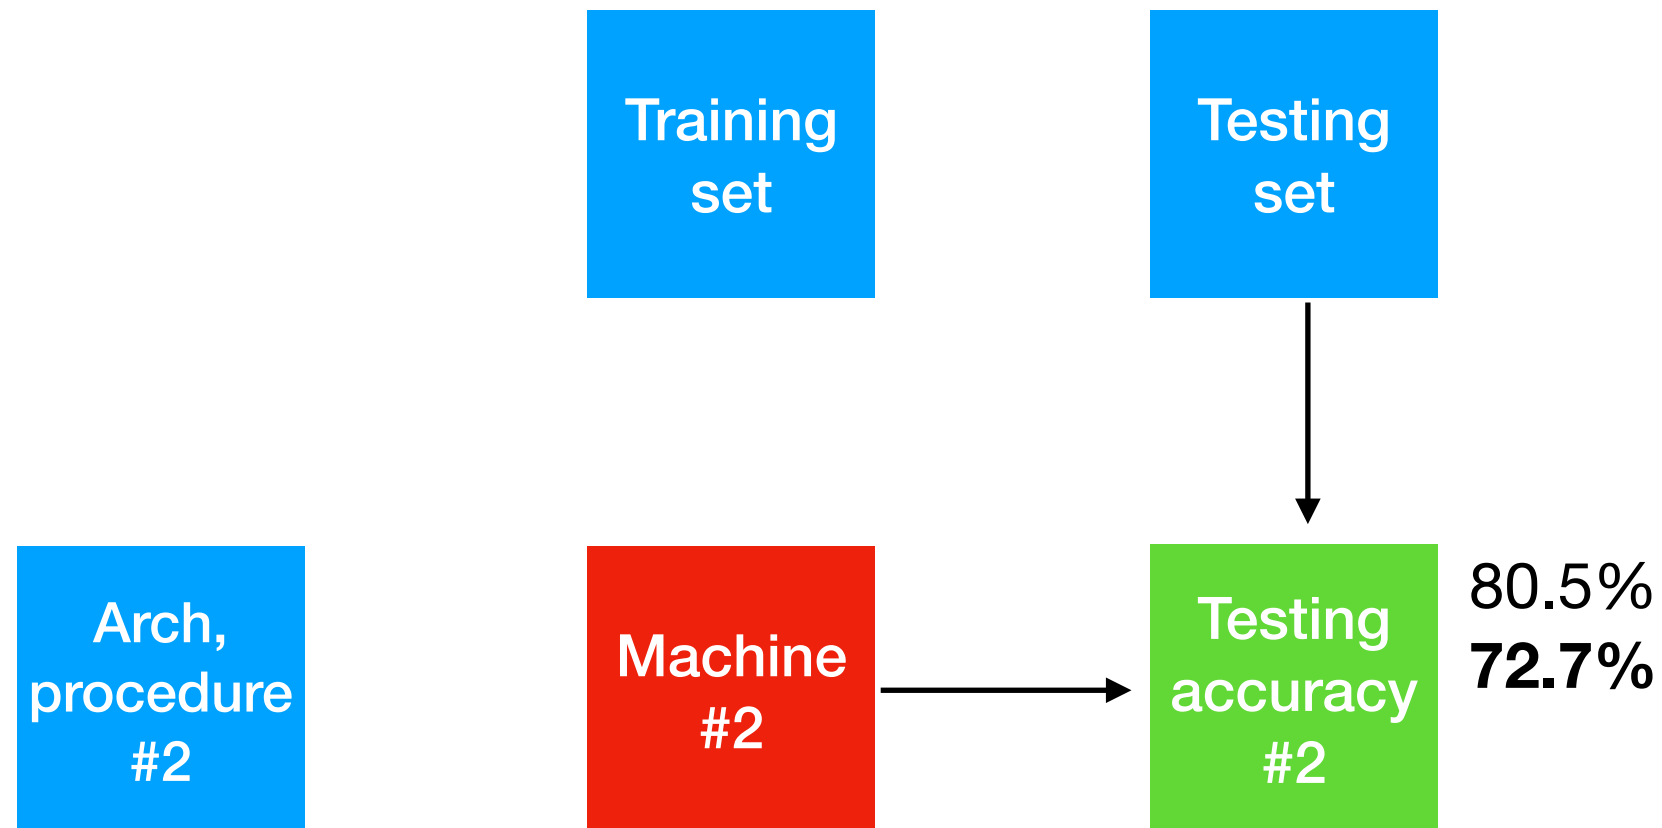
# Implicit cheating

- Choose a different design and try again!



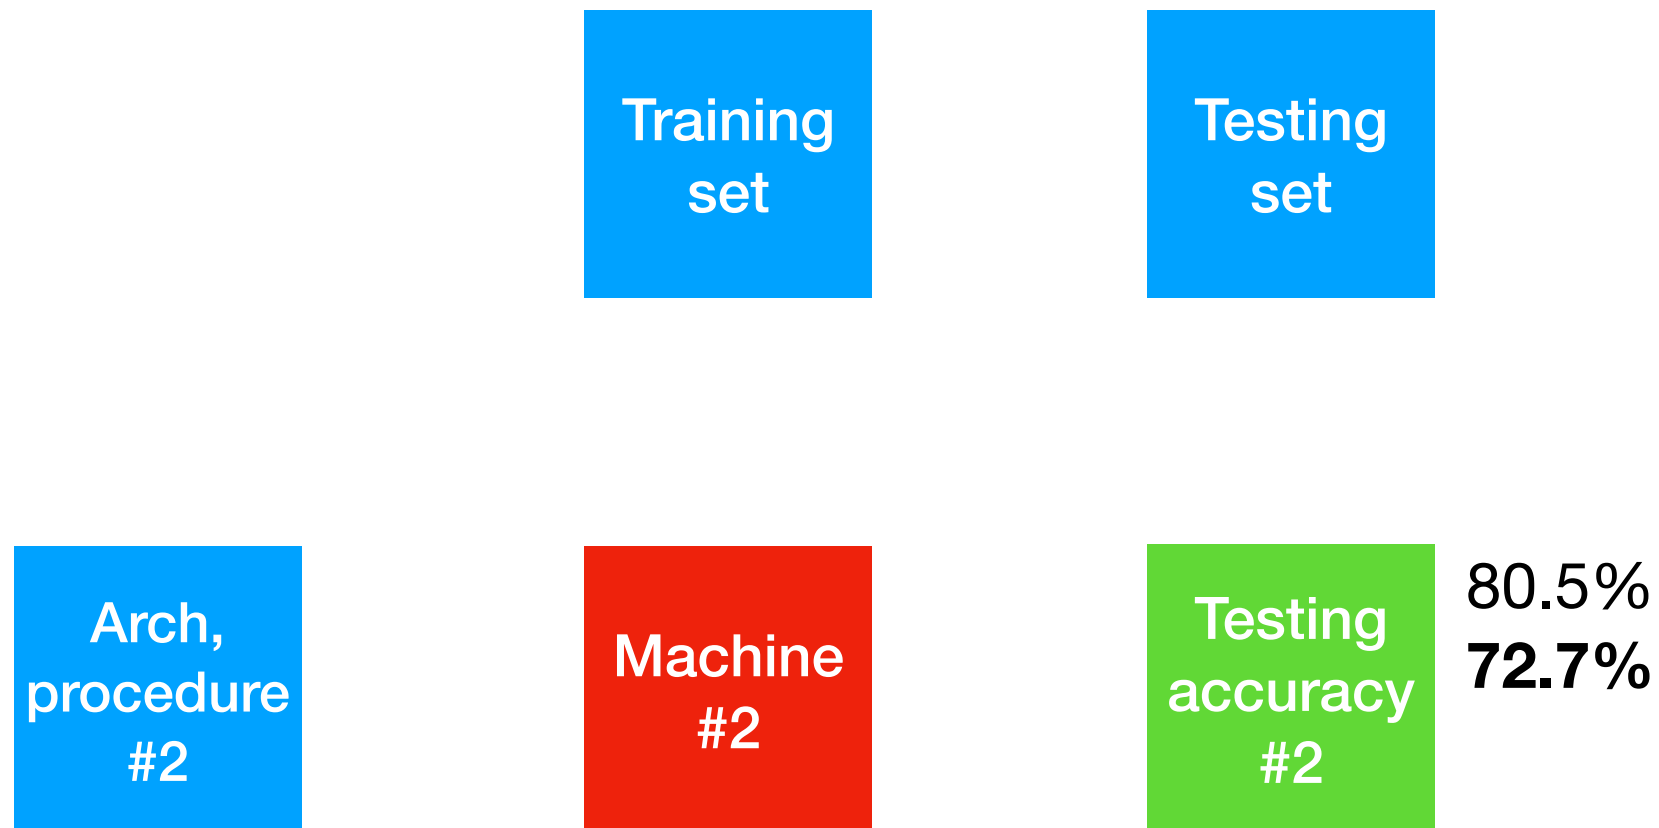
# Implicit cheating

- Evaluate the new machine on the testing set.



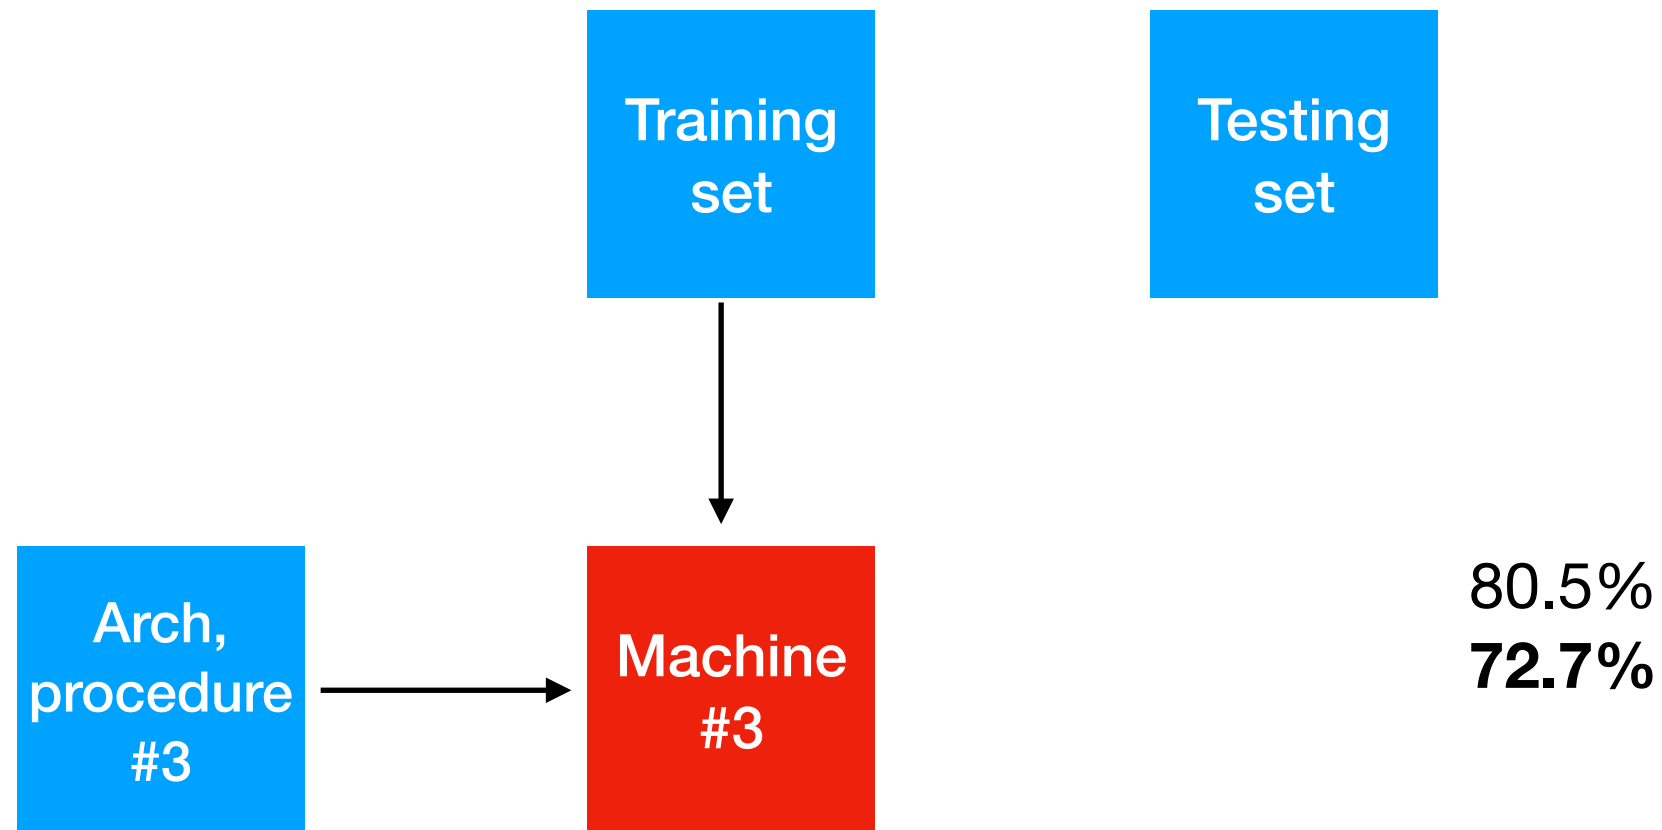
# Implicit cheating

- Accuracy still not good enough?



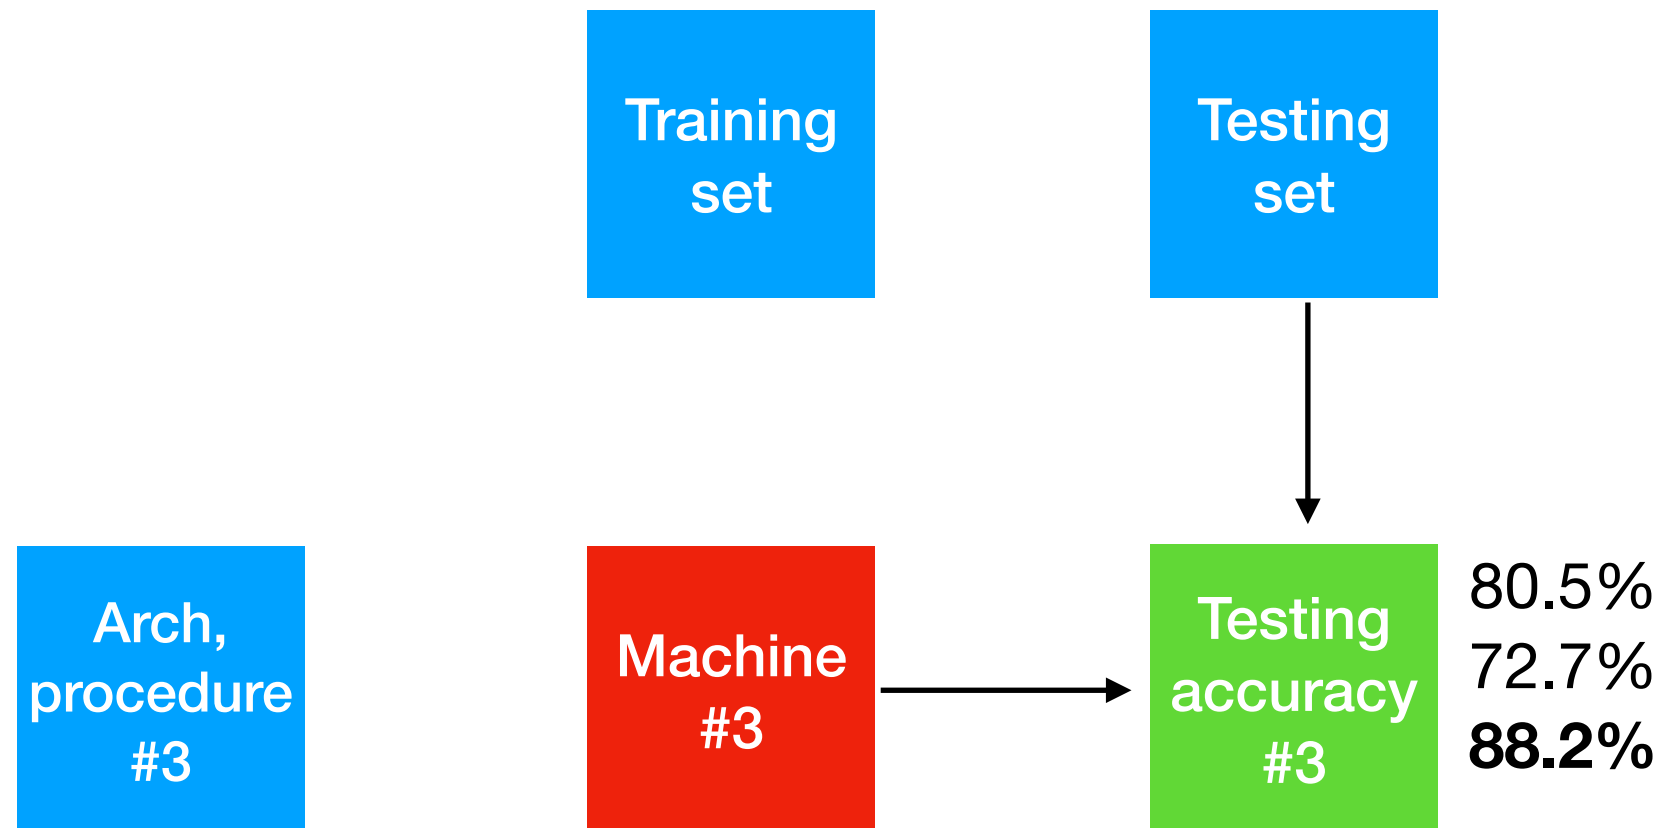
# Implicit cheating

- Choose yet another design and try again!



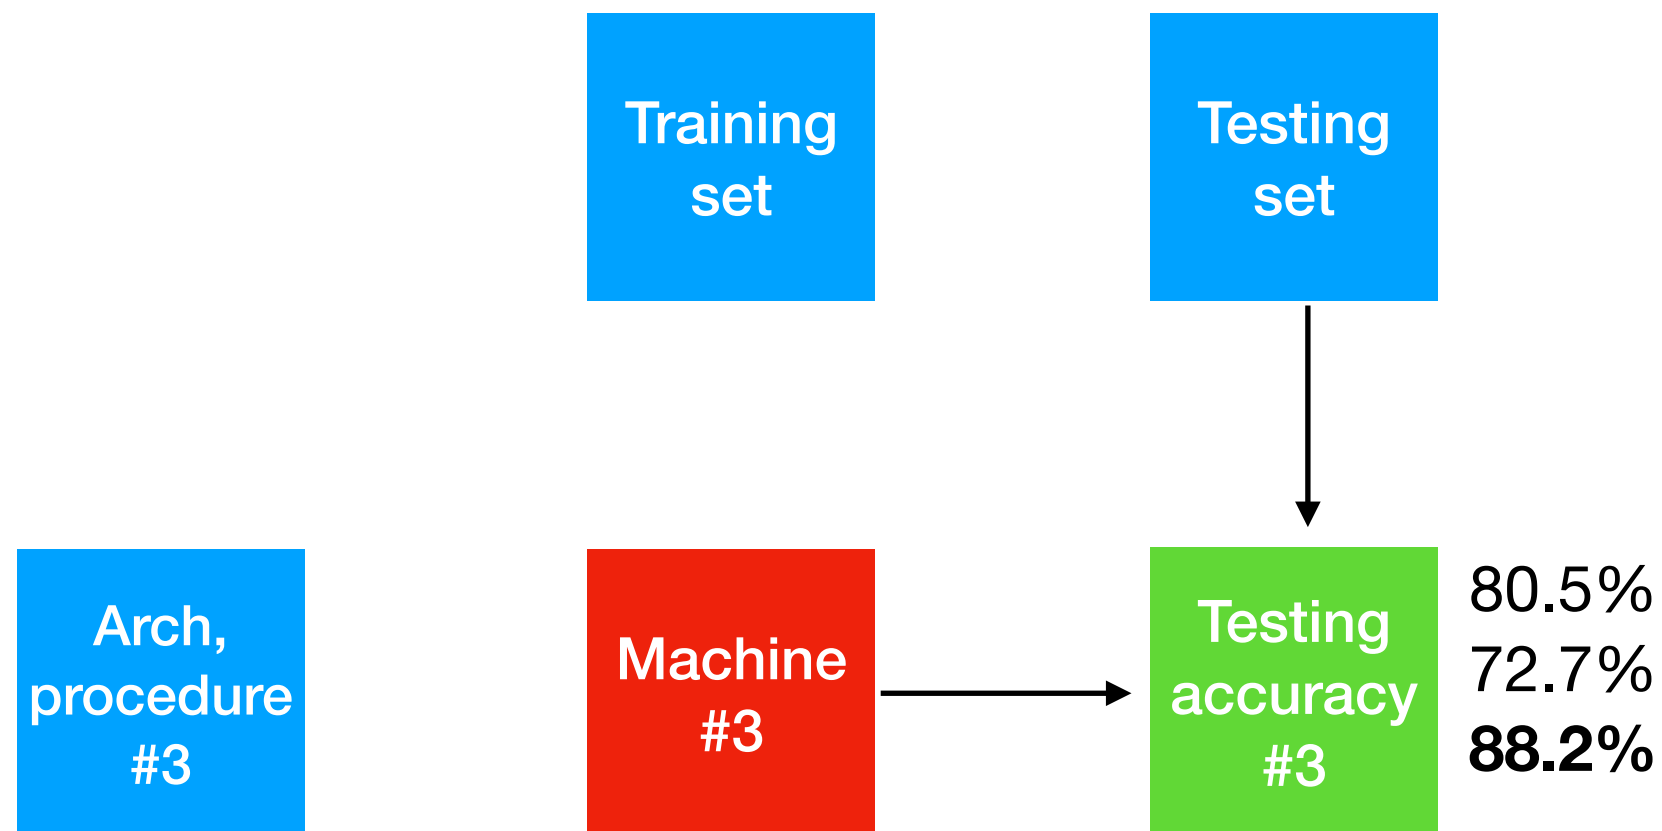
# Implicit cheating

- Evaluate the new machine on the testing set.



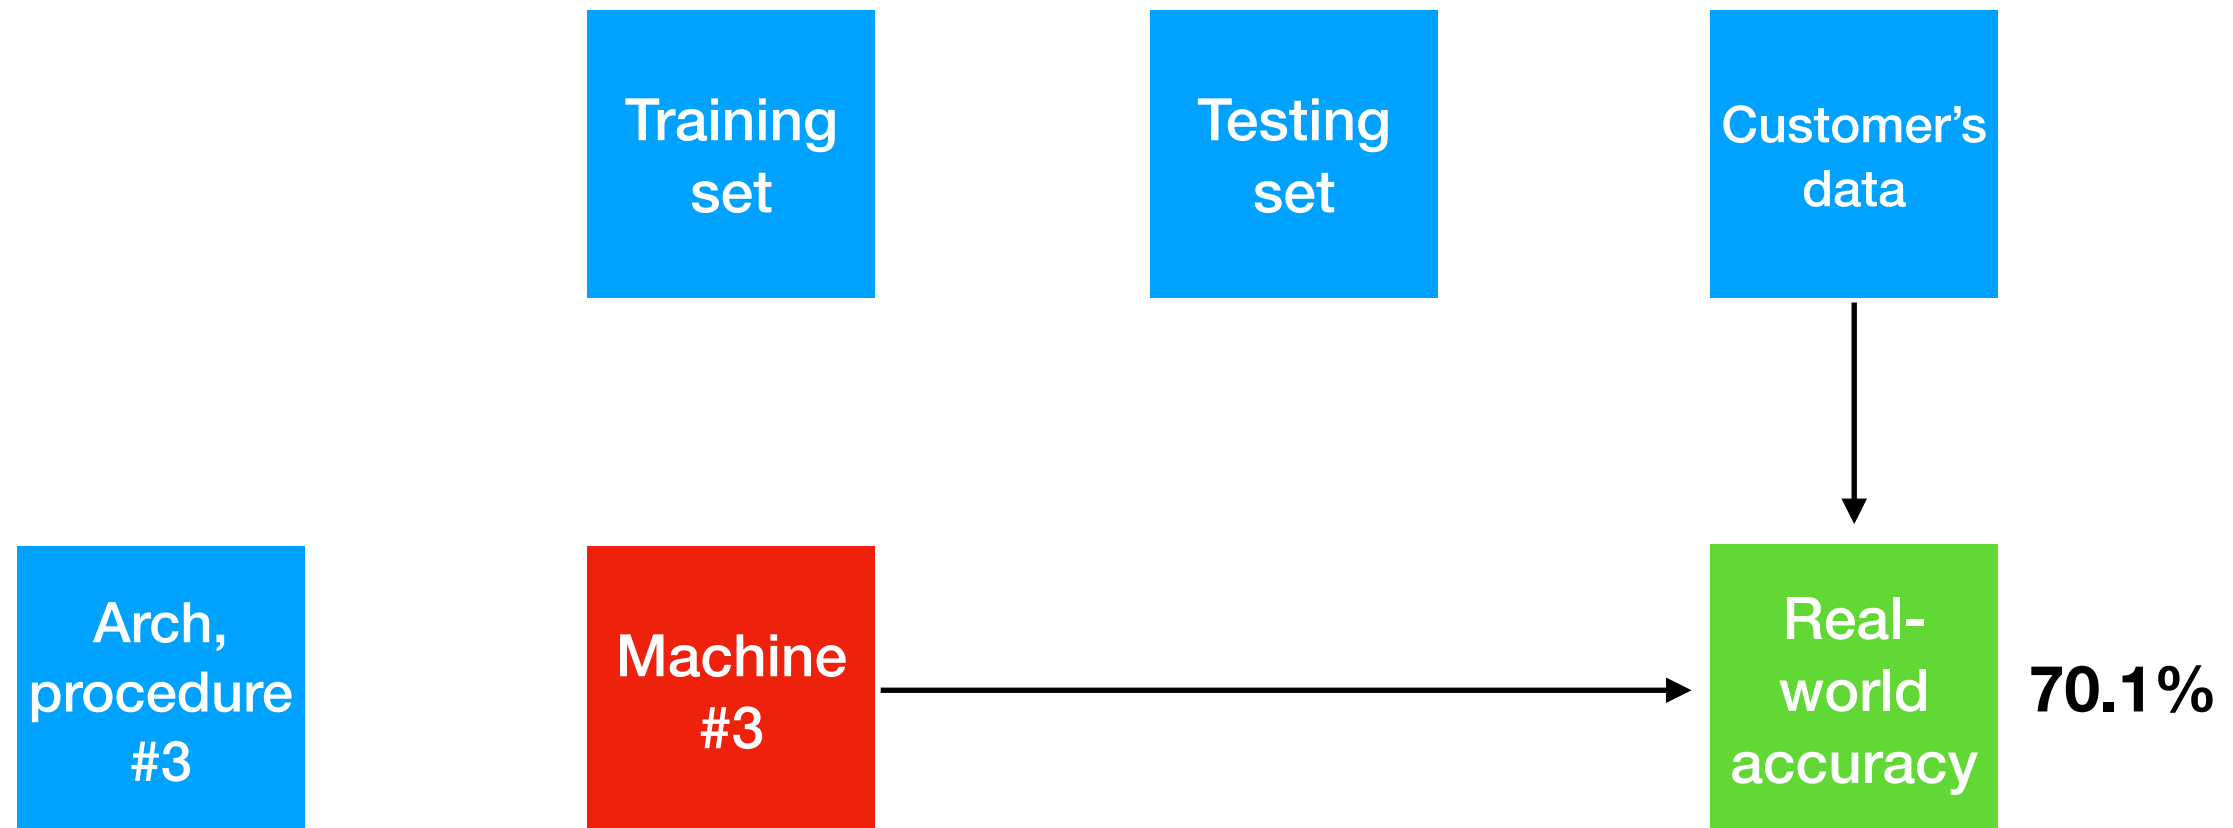
# Implicit cheating

- Much better! Let's keep machine #3 and sell it on Amazon!



# Implicit cheating

- Oops — the real-world accuracy was much less than what we estimated on the test set!



# Linear regression



# Limitations of our feature set

- So far, the predictors we have considered are very simple:
  - Is pixel  $(r_1, c_1)$  brighter than pixel  $(r_2, c_2)$ ?
- We can't even express simple relationships such as:
  - “Pixel  $(r_1, c_1)$  is at least 5 bigger than pixel  $(r_2, c_2)$ ”
  - “2 times pixel  $(r_1, c_1)$  is bigger than pixel  $(r_2, c_2)$ ”
  - “2 times pixel  $(r_1, c_1)$  plus 4 times pixel  $(r_2, c_2)$  is larger than pixel  $(r_3, c_3)$ ”.

# Linear regression

- We can harness these kinds of relationships using **linear regression**.
- Let's switch back to the age estimation problem...

# Linear algebra

- A **column vector** is a  $(n \times 1)$  matrix.
- A **row vector** is a  $(1 \times n)$  matrix.
- The **transpose** of  $(n \times k)$  matrix **A**, denoted **A**<sup>T</sup>, is  $(k \times n)$ .
- Multiplication of matrices **A** and **B**:
  - Only possible when: **A** is  $(n \times k)$  and **B** is  $(k \times m)$
  - Result:  $(n \times m)$

# Linear algebra

- The **inner product** between two column vectors (same length)  $\mathbf{x}$ ,  $\mathbf{y}$  can be written as:  $\mathbf{x}^T \mathbf{y}$

$$\begin{bmatrix} x_1 & \dots & x_m \end{bmatrix} \begin{bmatrix} y_1 \\ \vdots \\ y_m \end{bmatrix} = \sum_{i=1}^m x_i y_i$$

- An inner product produces a **scalar**.

# Linear algebra

- The **outer product** between two column vectors (same length)  $\mathbf{x}$ ,  $\mathbf{y}$  can be written as:  $\mathbf{xy}^T$ :

$$\begin{bmatrix} x_1 \\ \vdots \\ x_m \end{bmatrix} \begin{bmatrix} y_1 & \dots & y_m \end{bmatrix} = \begin{bmatrix} x_1 y_1 & \dots & x_1 y_m \\ \vdots & \ddots & \vdots \\ x_m y_1 & \dots & x_m y_m \end{bmatrix}$$

- An outer product produces a **matrix**.

# Example

$$\begin{bmatrix} 1 \\ 2 \end{bmatrix} \begin{bmatrix} 4 & -1 \end{bmatrix} = \begin{bmatrix} 4 & -1 \\ 8 & -2 \end{bmatrix}$$
$$\begin{bmatrix} -1 \\ 3 \end{bmatrix} \begin{bmatrix} 2 & 0 \end{bmatrix} = \begin{bmatrix} -2 & 0 \\ 6 & 0 \end{bmatrix}$$

# Linear algebra

- The **sum of multiple outer products** can be expressed as the multiplication of two matrices:

$$\begin{aligned}\mathbf{x}^{(1)}\mathbf{y}^{(1)\top} + \dots \mathbf{x}^{(n)}\mathbf{y}^{(n)\top} &= \sum_{i=1}^n \mathbf{x}^{(i)}\mathbf{y}^{(i)\top} \\ &= \begin{bmatrix} \left| \mathbf{x}^{(1)} \right| & \dots & \left| \mathbf{x}^{(n)} \right| \end{bmatrix} \begin{bmatrix} \text{---} & \mathbf{y}^{(1)} & \text{---} \\ & \vdots & \\ \text{---} & \mathbf{y}^{(n)} & \text{---} \end{bmatrix} \\ &\doteq \mathbf{XY}^\top\end{aligned}$$

# Example

$$\mathbf{x} \quad \mathbf{y}^T \quad \mathbf{xy}^T$$

$$\begin{bmatrix} 1 \\ 2 \end{bmatrix} \begin{bmatrix} 4 & -1 \end{bmatrix} = \begin{bmatrix} 4 & -1 \\ 8 & -2 \end{bmatrix}$$
$$\begin{bmatrix} -1 \\ 3 \end{bmatrix} \begin{bmatrix} 2 & 0 \end{bmatrix} = \begin{bmatrix} -2 & 0 \\ 6 & 0 \end{bmatrix}$$

$$\mathbf{X} \quad \mathbf{Y}^T \quad \mathbf{XY}^T$$

$$\begin{bmatrix} 1 & -1 \\ 2 & 3 \end{bmatrix} \begin{bmatrix} 4 & -1 \\ 2 & 0 \end{bmatrix} = \begin{bmatrix} 2 & -1 \\ 14 & -2 \end{bmatrix}$$



# Linear algebra

- Here's a special case:

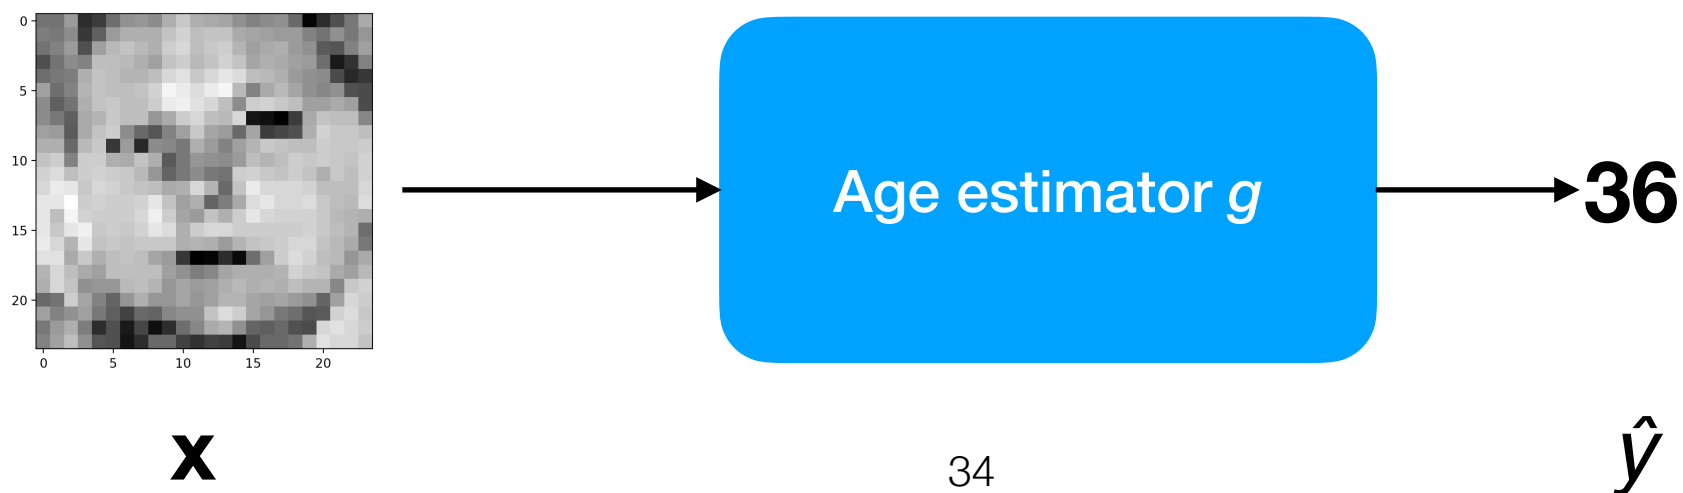
$$\begin{aligned}
 \mathbf{x}^{(1)}\mathbf{x}^{(1)\top} + \dots \mathbf{x}^{(n)}\mathbf{x}^{(n)\top} &= \sum_{i=1}^n \mathbf{x}^{(i)}\mathbf{x}^{(i)\top} \\
 &= \begin{bmatrix} \left| \mathbf{x}^{(1)} \right| & \dots & \left| \mathbf{x}^{(n)} \right| \end{bmatrix} \begin{bmatrix} \text{---} & \mathbf{x}^{(1)} & \text{---} \\ & \vdots & \\ \text{---} & \mathbf{x}^{(n)} & \text{---} \end{bmatrix} \\
 &\doteq \mathbf{X}\mathbf{X}^\top
 \end{aligned}$$

# Linear regression

- Linear regression is built as a linear combination of all the inputs  $\mathbf{x}$ :

$$\hat{y} = g(\mathbf{x}; \mathbf{w}) = \sum_{j=1}^m \text{image pixels } \mathbf{x}_j \mathbf{w}_j = \mathbf{x}^\top \mathbf{w}$$

- Here, we treat the image  $\mathbf{x}$  as a *vector* (even though it represents a 2-d image).

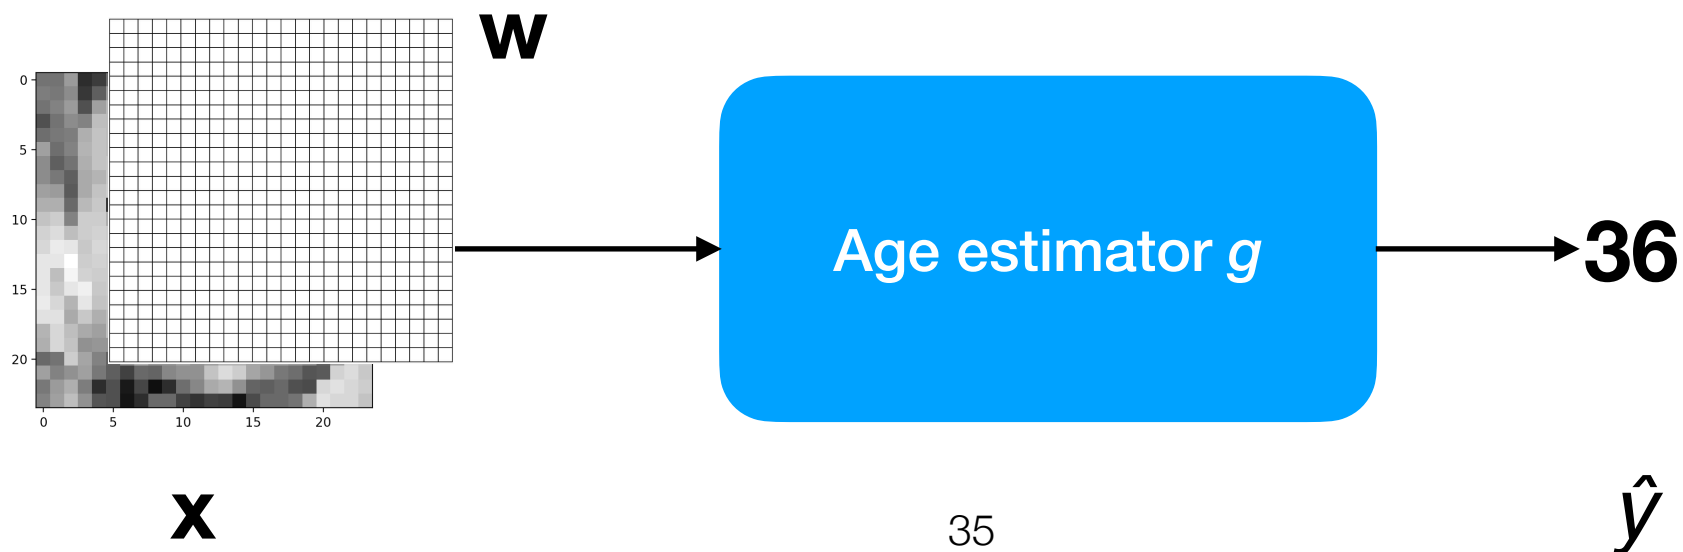


# Linear regression

- Linear regression is built as a linear combination of all the inputs  $\mathbf{x}$ :

$$\hat{y} = g(\mathbf{x}; \mathbf{w}) = \sum_{j=1}^m \text{image pixels } \mathbf{x}_j \mathbf{w}_j = \mathbf{x}^\top \mathbf{w}$$

- Vector  $\mathbf{w}$  represent an “overlay image” that weights the different pixel intensities of  $\mathbf{x}$ .



# Linear regression

- Imagine a 2x2 pixel “image”  $\mathbf{x}$  and a weight matrix  $\mathbf{w}$ :

2	5
0	3

$\mathbf{x}$

1	3
2	4

$\mathbf{w}$

- Then  $\hat{y} =$  ?

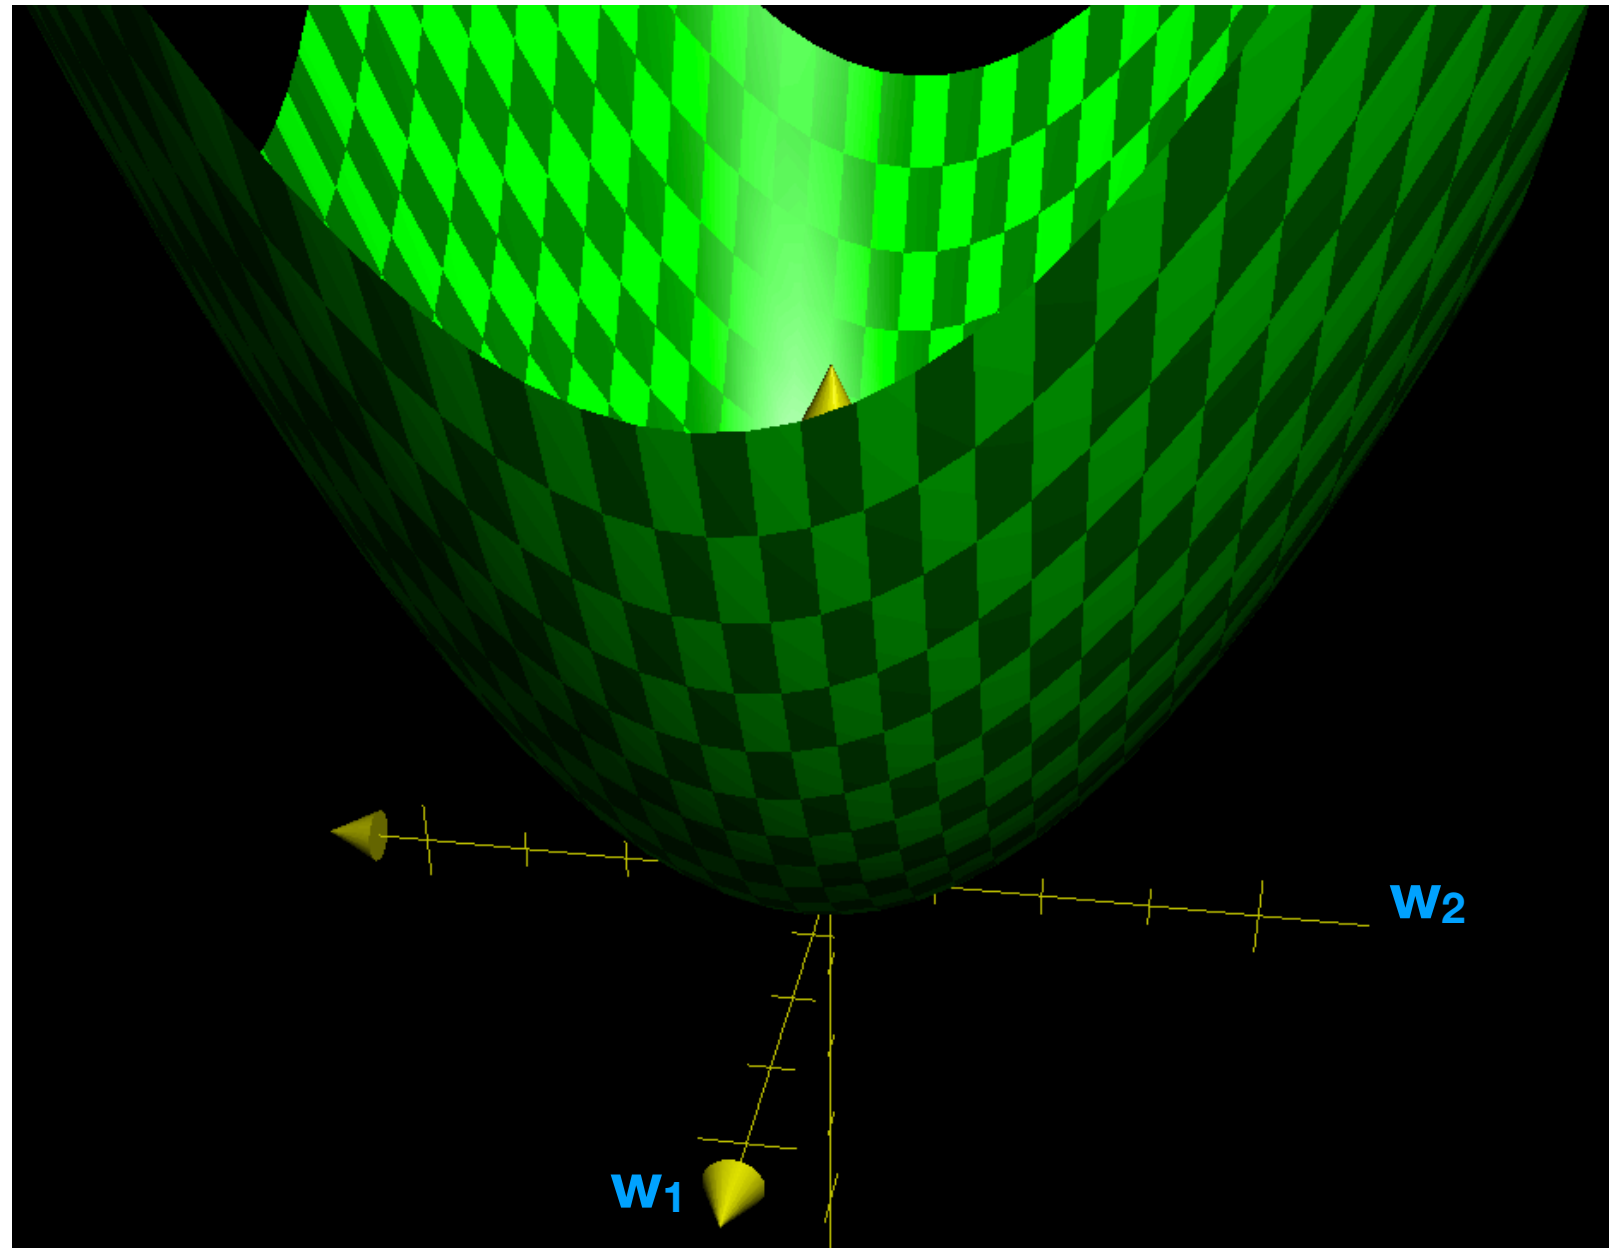
# Linear regression

- How should we choose each “weight”  $\mathbf{w}_j$ ?
- Let’s define the **loss** function that we seek to minimize:

$$\begin{aligned} f_{\text{MSE}}(\mathbf{y}, \hat{\mathbf{y}}; \mathbf{w}) &= \frac{1}{2n} \sum_{i=1}^n \left( g(\mathbf{x}^{(i)}; \mathbf{w}) - y^{(i)} \right)^2 \\ &= \frac{1}{2n} \sum_{i=1}^n \left( \mathbf{x}^{(i)\top} \mathbf{w} - y^{(i)} \right)^2 \end{aligned}$$

The 2 in the denominator will  
slightly simplify the algebra later...

# What does $f(w)$ look like?



# Linear regression

- $\mathbf{w}$  can be any real-valued vector; hence, we can use differential calculus to find the minimum of  $f_{\text{MSE}}$ .
- Just derive the **gradient** of  $f_{\text{MSE}}$  w.r.t.  $\mathbf{w}$ , set to 0, and solve.
- Since  $f_{\text{MSE}}$  is a convex function, we are guaranteed that this critical point is a global minimum.

# Gradient vector

- For a real-valued function  $f : \mathbb{R}^m \rightarrow \mathbb{R}$ , we define the gradient w.r.t.  $\mathbf{w}$  as:

$$\nabla_{\mathbf{w}} f = \begin{bmatrix} \frac{\partial f}{\partial \mathbf{w}_1} \\ \vdots \\ \frac{\partial f}{\partial \mathbf{w}_m} \end{bmatrix}$$

- In other words, the gradient is a column vector containing all first partial derivatives w.r.t.  $\mathbf{w}$ .



# Gradient vector: exercise 1

$$f(\mathbf{w}) = f\left(\begin{bmatrix} w_1 \\ w_2 \end{bmatrix}\right) = 3w_1^2 - \sin(2w_2)$$

$$\nabla_{\mathbf{w}} f(\mathbf{w}) = \begin{bmatrix} \phantom{0} \\ \phantom{0} \end{bmatrix}$$

# Gradient vector: exercise 2

$$f(\mathbf{x}, \mathbf{w}, y) = \frac{1}{2}(\mathbf{x}^\top \mathbf{w} - y)^2$$

$$\nabla_{\mathbf{w}} f(\mathbf{x}, \mathbf{w}, y) =$$

# Gradient vector: exercise 2

$$\begin{aligned} f(\mathbf{x}, \mathbf{w}, y) &= \frac{1}{2}(\mathbf{x}^\top \mathbf{w} - y)^2 \\ &= \frac{1}{2}(x_1 w_1 + x_2 w_2 - y)^2 \end{aligned}$$

$$\nabla_{\mathbf{w}} f(\mathbf{x}, \mathbf{w}, y) =$$

# Solving for $\mathbf{w}$

- The gradient of  $f_{\text{MSE}}$  is thus:

$$\begin{aligned}\nabla_{\mathbf{w}} f_{\text{MSE}}(\mathbf{y}, \hat{\mathbf{y}}; \mathbf{w}) &= \nabla_{\mathbf{w}} \left[ \frac{1}{2n} \sum_{i=1}^n \left( \mathbf{x}^{(i)\top} \mathbf{w} - y^{(i)} \right)^2 \right] \\ &= \frac{1}{2n} \sum_{i=1}^n \nabla_{\mathbf{w}} \left[ \left( \mathbf{x}^{(i)\top} \mathbf{w} - y^{(i)} \right)^2 \right] \\ &= \frac{1}{n} \sum_{i=1}^n \mathbf{x}^{(i)} \left( \mathbf{x}^{(i)\top} \mathbf{w} - y^{(i)} \right)\end{aligned}$$

# Solving for $\mathbf{w}$

- By setting to 0, splitting the sum apart, and solving, we reach the solution:

$$\frac{1}{n} \sum_{i=1}^n \mathbf{x}^{(i)} \left( \mathbf{x}^{(i)\top} \mathbf{w} - y^{(i)} \right)$$

$$0 = \sum_i \mathbf{x}^{(i)} \mathbf{x}^{(i)\top} \mathbf{w} - \sum_i \mathbf{x}^{(i)} y^{(i)}$$

# Solving for $\mathbf{w}$

- By setting to 0, splitting the sum apart, and solving, we reach the solution:

$$\frac{1}{n} \sum_{i=1}^n \mathbf{x}^{(i)} \left( \mathbf{x}^{(i)\top} \mathbf{w} - y^{(i)} \right)$$

$$0 = \sum_i \mathbf{x}^{(i)} \mathbf{x}^{(i)\top} \mathbf{w} - \sum_i \mathbf{x}^{(i)} y^{(i)}$$

$$\sum_i \mathbf{x}^{(i)} \mathbf{x}^{(i)\top} \mathbf{w} = \sum_i \mathbf{x}^{(i)} y^{(i)}$$

# Solving for $\mathbf{w}$

- By setting to 0, splitting the sum apart, and solving, we reach the solution:

$$\frac{1}{n} \sum_{i=1}^n \mathbf{x}^{(i)} \left( \mathbf{x}^{(i)\top} \mathbf{w} - y^{(i)} \right)$$

$$0 = \sum_i \mathbf{x}^{(i)} \mathbf{x}^{(i)\top} \mathbf{w} - \sum_i \mathbf{x}^{(i)} y^{(i)}$$

$$\sum_i \mathbf{x}^{(i)} \mathbf{x}^{(i)\top} \mathbf{w} = \sum_i \mathbf{x}^{(i)} y^{(i)}$$

$$\mathbf{w} = \left( \sum_i \mathbf{x}^{(i)} \mathbf{x}^{(i)\top} \right)^{-1} \sum_i \mathbf{x}^{(i)} y^{(i)}$$

# Linear regression: matrix notation

- To compute  $\mathbf{w}$ , do *not* use `np.linalg.inv`.
- Instead, use `np.linalg.solve`, which avoids explicitly computing the matrix inverse.



# Linear regression: matrix notation

- Let's define a matrix  $\mathbf{X}$  to contain all the training images:

$$\mathbf{X} = \begin{bmatrix} \mathbf{x}^{(1)} & \dots & \mathbf{x}^{(n)} \end{bmatrix}$$

- In statistics,  $\mathbf{X}$  is called the **design matrix**.
- Let's define vector  $\mathbf{y}$  to contain all the training labels:

$$\mathbf{y} = \begin{bmatrix} y^{(1)} \\ \vdots \\ y^{(n)} \end{bmatrix}$$

# Linear regression: matrix notation

- Using summation notation, we derived:

$$\mathbf{w} = \left( \sum_{i=1}^n \mathbf{x}^{(i)} \mathbf{x}^{(i)\top} \right)^{-1} \left( \sum_{i=1}^n \mathbf{x}^{(i)} y^{(i)} \right)$$

- Using matrix notation, we can write the solution as:

$\mathbf{w} =$

?

# Linear regression: matrix notation

- Once we've "trained" the weights  $\mathbf{w}$ , we can estimate the  $y$ -value (label) for any  $\mathbf{x}$ .
- We can compute the  $\{ \hat{y}^{(i)} \}$  for a set of images  $\{ \mathbf{x}^{(i)} \}$  in one-fell-swoop using matrix operations.
- Let's define our design matrix  $\mathbf{X}$  as before:

$$\mathbf{X} = \begin{bmatrix} | & & | \\ \mathbf{x}^{(1)} & \dots & \mathbf{x}^{(n)} \\ | & & | \end{bmatrix}$$

- Then our estimates of the labels is given by:

$$\hat{\mathbf{y}} = \mathbf{X}^\top \mathbf{w}$$

# Linear regression: matrix notation

- Suppose we have  $n$  images, each with just 2 pixels.

$$\hat{y} = \mathbf{X}^\top \mathbf{w}$$

# Linear regression: matrix notation

- Suppose we have  $n$  images, each with just 2 pixels.

$$\begin{aligned}\hat{y} &= \mathbf{X}^\top \mathbf{w} \\ &= \begin{bmatrix} \mathbf{x}_1^{(1)} & \dots & \mathbf{x}_1^{(n)} \\ \mathbf{x}_2^{(1)} & \dots & \mathbf{x}_2^{(n)} \end{bmatrix}^\top \begin{bmatrix} w_1 \\ w_2 \end{bmatrix}\end{aligned}$$

This is the index of  
the *image*.

# Linear regression: matrix notation

- Suppose we have  $n$  images, each with just 2 pixels.

$$\begin{aligned}\hat{y} &= \mathbf{X}^\top \mathbf{w} \\ &= \begin{bmatrix} \mathbf{x}_1^{(1)} & \dots & \mathbf{x}_1^{(n)} \\ \mathbf{x}_2^{(1)} & \dots & \mathbf{x}_2^{(n)} \end{bmatrix}^\top \begin{bmatrix} w_1 \\ w_2 \end{bmatrix}\end{aligned}$$

This is the index of  
the *pixel*.

# Linear regression: matrix notation

- Suppose we have  $n$  images, each with just 2 pixels.

$$\begin{aligned}\hat{y} &= \mathbf{X}^\top \mathbf{w} \\ &= \begin{bmatrix} \mathbf{x}_1^{(1)} & \dots & \mathbf{x}_1^{(n)} \\ \mathbf{x}_2^{(1)} & \dots & \mathbf{x}_2^{(n)} \end{bmatrix}^\top \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} \\ &= \begin{bmatrix} \mathbf{x}_1^{(1)} & \mathbf{x}_2^{(1)} \\ \vdots & \vdots \\ \mathbf{x}_1^{(n)} & \mathbf{x}_2^{(n)} \end{bmatrix} \begin{bmatrix} w_1 \\ w_2 \end{bmatrix}\end{aligned}$$

# Linear regression: matrix notation

- Suppose we have  $n$  images, each with just 2 pixels.

$$\begin{aligned}\hat{y} &= \mathbf{X}^\top \mathbf{w} \\ &= \begin{bmatrix} \mathbf{x}_1^{(1)} & \dots & \mathbf{x}_1^{(n)} \\ \mathbf{x}_2^{(1)} & \dots & \mathbf{x}_2^{(n)} \end{bmatrix}^\top \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} \\ &= \begin{bmatrix} \mathbf{x}_1^{(1)} & \mathbf{x}_2^{(1)} \\ \vdots & \vdots \\ \mathbf{x}_1^{(n)} & \mathbf{x}_2^{(n)} \end{bmatrix} \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} \\ &= \begin{bmatrix} \mathbf{x}_1^{(1)} w_1 + \mathbf{x}_2^{(1)} w_2 \\ \vdots \\ \mathbf{x}_1^{(n)} w_1 + \mathbf{x}_2^{(n)} w_2 \end{bmatrix}\end{aligned}$$