

CS 4342: Class 4

Jacob Whitehill

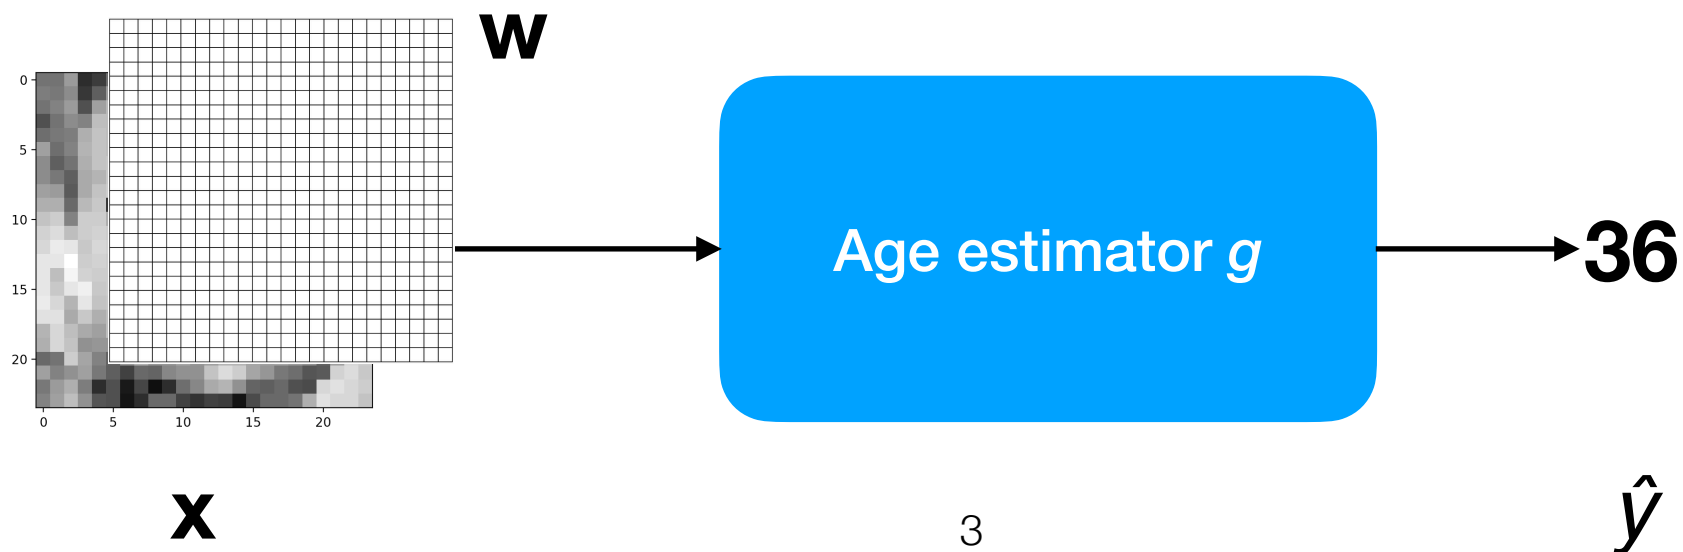
Linear regression

Linear regression

- Linear regression is built as a linear combination of all the inputs \mathbf{x} :

$$\hat{y} = g(\mathbf{x}; \mathbf{w}) = \sum_{j=1}^m \text{image pixels } \mathbf{x}_j \mathbf{w}_j = \mathbf{x}^\top \mathbf{w}$$

- Vector \mathbf{w} represent an “overlay image” that weights the different pixel intensities of \mathbf{x} .



Solving for \mathbf{w}

- The gradient of f_{MSE} is thus:

$$\begin{aligned}\nabla_{\mathbf{w}} f_{\text{MSE}}(\mathbf{y}, \hat{\mathbf{y}}; \mathbf{w}) &= \nabla_{\mathbf{w}} \left[\frac{1}{2n} \sum_{i=1}^n \left(\mathbf{x}^{(i)\top} \mathbf{w} - y^{(i)} \right)^2 \right] \\ &= \frac{1}{2n} \sum_{i=1}^n \nabla_{\mathbf{w}} \left[\left(\mathbf{x}^{(i)\top} \mathbf{w} - y^{(i)} \right)^2 \right] \\ &= \frac{1}{n} \sum_{i=1}^n \mathbf{x}^{(i)} \left(\mathbf{x}^{(i)\top} \mathbf{w} - y^{(i)} \right)\end{aligned}$$

Solving for \mathbf{w}

- By setting to 0, splitting the sum apart, and solving, we reach the solution:

$$\frac{1}{n} \sum_{i=1}^n \mathbf{x}^{(i)} \left(\mathbf{x}^{(i)\top} \mathbf{w} - y^{(i)} \right)$$

$$0 = \sum_i \mathbf{x}^{(i)} \mathbf{x}^{(i)\top} \mathbf{w} - \sum_i \mathbf{x}^{(i)} y^{(i)}$$

$$\sum_i \mathbf{x}^{(i)} \mathbf{x}^{(i)\top} \mathbf{w} = \sum_i \mathbf{x}^{(i)} y^{(i)}$$

$$\mathbf{w} = \left(\sum_i \mathbf{x}^{(i)} \mathbf{x}^{(i)\top} \right)^{-1} \sum_i \mathbf{x}^{(i)} y^{(i)}$$

Linear regression: matrix notation

- To compute \mathbf{w} , do *not* use `np.linalg.inv`.
- Instead, use `np.linalg.solve`, which avoids explicitly computing the matrix inverse.
- Show `age_demo.py`.

Linear regression: matrix notation

- Let's define a matrix \mathbf{X} to contain all the training images:

$$\mathbf{X} = \begin{bmatrix} \mathbf{x}^{(1)} & \dots & \mathbf{x}^{(n)} \end{bmatrix}$$

- In statistics, \mathbf{X} is called the **design matrix**.
- Let's define vector \mathbf{y} to contain all the training labels:

$$\mathbf{y} = \begin{bmatrix} y^{(1)} \\ \vdots \\ y^{(n)} \end{bmatrix}$$

Linear regression: matrix notation

- Using summation notation, we derived:

$$\mathbf{w} = \left(\sum_{i=1}^n \mathbf{x}^{(i)} \mathbf{x}^{(i)\top} \right)^{-1} \left(\sum_{i=1}^n \mathbf{x}^{(i)} y^{(i)} \right)$$

- Using matrix notation, we can write the solution as:

$$\mathbf{w} =$$



Linear regression: matrix notation

- Using summation notation, we derived:

$$\mathbf{w} = \left(\sum_{i=1}^n \mathbf{x}^{(i)} \mathbf{x}^{(i)\top} \right)^{-1} \left(\sum_{i=1}^n \mathbf{x}^{(i)} y^{(i)} \right)$$

- Using matrix notation, we can write the solution as:

$$\mathbf{w} = (\mathbf{X}\mathbf{X}^\top)^{-1} \mathbf{X}\mathbf{y}$$

Linear regression: matrix notation

- Once we've "trained" the weights \mathbf{w} , we can estimate the y -value (label) for any \mathbf{x} .
- We can compute the $\{ \hat{y}^{(i)} \}$ for a set of images $\{ \mathbf{x}^{(i)} \}$ in one-fell-swoop using matrix operations.
- Let's define our design matrix \mathbf{X} as before:

$$\mathbf{X} = \begin{bmatrix} | & & | \\ \mathbf{x}^{(1)} & \dots & \mathbf{x}^{(n)} \\ | & & | \end{bmatrix}$$

- Then our estimates of the labels is given by:

$$\hat{\mathbf{y}} = \mathbf{X}^{\top} \mathbf{w}$$

Linear regression: matrix notation

- Suppose we have n images, each with just 2 pixels.

$$\hat{y} = \mathbf{X}^\top \mathbf{w}$$

Linear regression: matrix notation

- Suppose we have n images, each with just 2 pixels.

$$\begin{aligned}\hat{y} &= \mathbf{X}^\top \mathbf{w} \\ &= \begin{bmatrix} \mathbf{x}_1^{(1)} & \dots & \mathbf{x}_1^{(n)} \\ \mathbf{x}_2^{(1)} & \dots & \mathbf{x}_2^{(n)} \end{bmatrix}^\top \begin{bmatrix} w_1 \\ w_2 \end{bmatrix}\end{aligned}$$

This is the index of
the *image*.

Linear regression: matrix notation

- Suppose we have n images, each with just 2 pixels.

$$\begin{aligned}\hat{y} &= \mathbf{X}^\top \mathbf{w} \\ &= \begin{bmatrix} \mathbf{x}_1^{(1)} & \dots & \mathbf{x}_1^{(n)} \\ \mathbf{x}_2^{(1)} & \dots & \mathbf{x}_2^{(n)} \end{bmatrix}^\top \begin{bmatrix} w_1 \\ w_2 \end{bmatrix}\end{aligned}$$

This is the index of
the *pixel*.

Linear regression: matrix notation

- Suppose we have n images, each with just 2 pixels.

$$\begin{aligned}\hat{y} &= \mathbf{X}^\top \mathbf{w} \\ &= \begin{bmatrix} \mathbf{x}_1^{(1)} & \dots & \mathbf{x}_1^{(n)} \\ \mathbf{x}_2^{(1)} & \dots & \mathbf{x}_2^{(n)} \end{bmatrix}^\top \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} \\ &= \begin{bmatrix} \mathbf{x}_1^{(1)} & \mathbf{x}_2^{(1)} \\ \vdots & \vdots \\ \mathbf{x}_1^{(n)} & \mathbf{x}_2^{(n)} \end{bmatrix} \begin{bmatrix} w_1 \\ w_2 \end{bmatrix}\end{aligned}$$

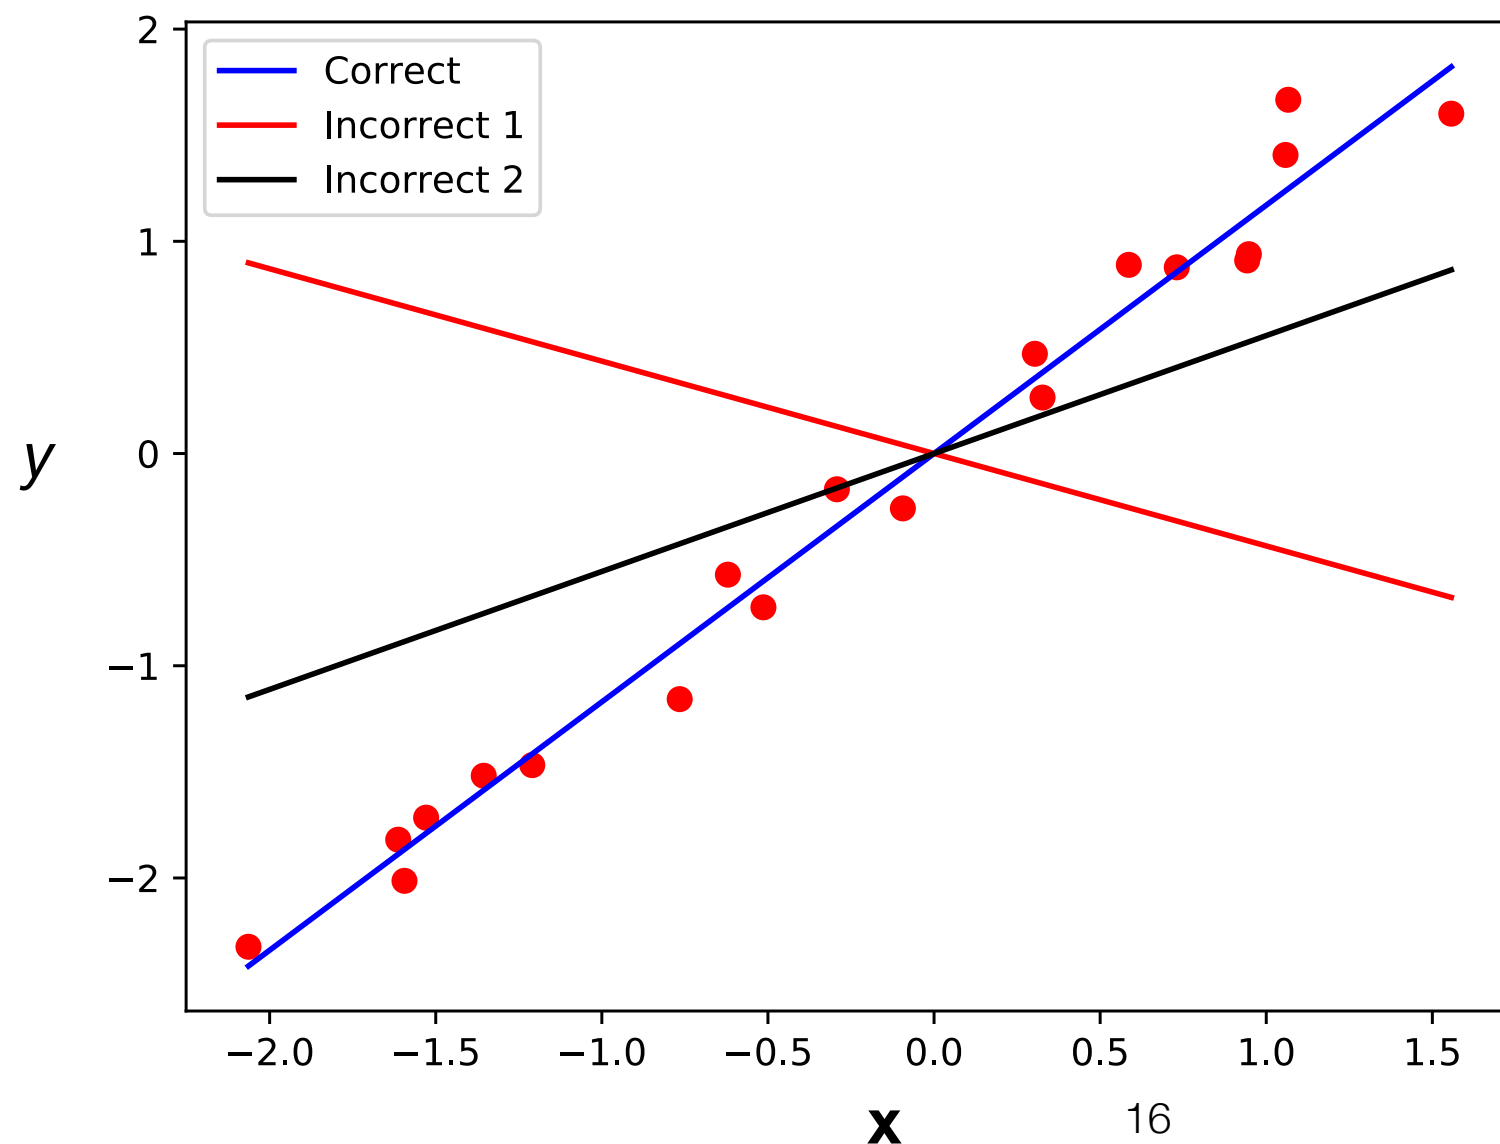
Linear regression: matrix notation

- Suppose we have n images, each with just 2 pixels.

$$\begin{aligned}\hat{y} &= \mathbf{X}^\top \mathbf{w} \\ &= \begin{bmatrix} \mathbf{x}_1^{(1)} & \dots & \mathbf{x}_1^{(n)} \\ \mathbf{x}_2^{(1)} & \dots & \mathbf{x}_2^{(n)} \end{bmatrix}^\top \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} \\ &= \begin{bmatrix} \mathbf{x}_1^{(1)} & \mathbf{x}_2^{(1)} \\ \vdots & \vdots \\ \mathbf{x}_1^{(n)} & \mathbf{x}_2^{(n)} \end{bmatrix} \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} \\ &= \begin{bmatrix} \mathbf{x}_1^{(1)} w_1 + \mathbf{x}_2^{(1)} w_2 \\ \vdots \\ \mathbf{x}_1^{(n)} w_1 + \mathbf{x}_2^{(n)} w_2 \end{bmatrix}\end{aligned}$$

1-d example

- Linear regression finds the weight vector \mathbf{w} that minimizes the f_{MSE} . Here's an example where each \mathbf{x} is just 1-d...



The best \mathbf{w} is the one such that $f_{\text{MSE}}(\mathbf{y}, \hat{\mathbf{y}})$ is as small as possible, where each $\hat{y} = \mathbf{x}^T \mathbf{w}$.

Bias term

- In order to account for target values y with non-zero mean, we could add a **bias term** b to our model:

$$\hat{y} = \mathbf{x}^\top \mathbf{w} + b$$

- We could then compute the gradient w.r.t. both \mathbf{w} and b and solve.

$$\begin{aligned}\nabla_{\mathbf{w}} f_{\text{MSE}}(\mathbf{y}, \hat{\mathbf{y}}; \mathbf{w}, b) &= \nabla_{\mathbf{w}} \left[\frac{1}{2n} \sum_{i=1}^n \left(\mathbf{x}^{(i)\top} \mathbf{w} + b - y^{(i)} \right)^2 \right] \\ \nabla_b f_{\text{MSE}}(\mathbf{y}, \hat{\mathbf{y}}; \mathbf{w}, b) &= \nabla_b \left[\frac{1}{2n} \sum_{i=1}^n \left(\mathbf{x}^{(i)\top} \mathbf{w} + b - y^{(i)} \right)^2 \right]\end{aligned}$$

Bias term

- Alternatively, we can implicitly include a bias term by augmenting each input vector \mathbf{x} with a 1 at the end:

$$\tilde{\mathbf{x}} = \begin{bmatrix} \mathbf{x} \\ 1 \end{bmatrix}$$

- Correspondingly, our weight vector \mathbf{w} will have an extra component (bias term) at the end.

$$\tilde{\mathbf{w}} = \begin{bmatrix} \mathbf{w} \\ b \end{bmatrix}$$

Bias term

- To see why, notice that:

$$\begin{aligned}\hat{y} &= \tilde{\mathbf{x}}^\top \tilde{\mathbf{w}} \\ &= \begin{bmatrix} \mathbf{x} \\ 1 \end{bmatrix}^\top \begin{bmatrix} \mathbf{w} \\ b \end{bmatrix} \\ &= \begin{bmatrix} \mathbf{x}^\top & 1 \end{bmatrix} \begin{bmatrix} \mathbf{w} \\ b \end{bmatrix} \\ &= \mathbf{x}^\top \mathbf{w} + b\end{aligned}$$

Bias term

- We can find the optimal \mathbf{w} and b based on all the training data using matrix notation.
- First define an augmented design matrix:

$$\tilde{\mathbf{X}} = \begin{bmatrix} \mathbf{x}^{(1)} & \dots & \mathbf{x}^{(n)} \\ 1 & \dots & 1 \end{bmatrix}$$

- Then compute:

$$\tilde{\mathbf{w}} = \left(\tilde{\mathbf{X}} \tilde{\mathbf{X}}^\top \right)^{-1} \tilde{\mathbf{X}} \mathbf{y}$$

Fairness in ML

Fairness in ML

- Consider the following definition of ML fairness:
 - The machine's accuracy should be equal across all demographic subgroups on which it is tested.

Exercise

- Suppose we have trained a classifier to perceive whether a person is smiling based on their face image, and suppose its test accuracy (PC) on male & female faces is:
 - Male: 42%; female: 55%
(You may assume that people from both genders smile with 50% probability.)
- Describe 3 possible reasons for why the test accuracy may differ between male and female faces.

Iterative solution to linear regression

Linear regression

- Linear regression is one of the few ML algorithms that has an analytical solution:

$$\mathbf{w} = (\mathbf{X}\mathbf{X}^\top)^{-1} \mathbf{X}\mathbf{y}$$

- **Analytical solution:** there is a closed formula for the answer.

Linear regression


- Alternatively, linear regression can be solved numerically using gradient descent.
- **Numerical solution:** need to iterate (according to some algorithm) many times to *approximate* the optimal value.
- Gradient descent is more laborious to code than the one-shot solution, but it generalizes to a wide variety of ML models.

Gradient descent

- Gradient descent is a **hill climbing algorithm** that uses the gradient (aka slope) to decide which way to “move” \mathbf{w} to reduce the objective function (e.g., f_{MSE}).


Gradient descent

- Suppose we just guess an initial value for w (e.g., -2.1).
- How can we make it better — increase it or decrease it?



Gradient descent


- Suppose we just guess an initial value for w (e.g., -2.1).
- How can we make it better — **increase** it or decrease it?
- What does the **slope** of f_{MSE} tell us to do?



The slope at $f_{\text{MSE}}(-2.1)$ is *negative*, i.e., we can *decrease* our cost by *increasing* w .

Gradient descent

- Or maybe our initial guess for w was 3.9.
- How can we make it better — increase it or **decrease** it?
- What does the **slope** of f_{MSE} tell us to do?



The slope at $f_{\text{MSE}}(3.9)$ is *positive*,
i.e., we can *decrease* our cost
by *decreasing* w .

Gradient descent

- How do we know the slope? Compute the **gradient** of f_{MSE} w.r.t. \mathbf{w} :

$$\begin{aligned}\nabla_{\mathbf{w}} f_{\text{MSE}}(\mathbf{y}, \hat{\mathbf{y}}; \mathbf{w}) &= \nabla_{\mathbf{w}} \left[\frac{1}{2n} \sum_{i=1}^n \left(\mathbf{x}^{(i)\top} \mathbf{w} - y^{(i)} \right)^2 \right] \\ &= \frac{1}{2n} \sum_{i=1}^n \nabla_{\mathbf{w}} \left[\left(\mathbf{x}^{(i)\top} \mathbf{w} - y^{(i)} \right)^2 \right] \\ &= \frac{1}{n} \sum_{i=1}^n \mathbf{x}^{(i)} \left(\mathbf{x}^{(i)\top} \mathbf{w} - y^{(i)} \right) \\ &= \frac{1}{n} \mathbf{X} (\mathbf{X}^\top \mathbf{w} - \mathbf{y})\end{aligned}$$

Gradient descent


- How do we know the slope? Compute the **gradient** of f_{MSE} w.r.t. \mathbf{w} :

$$\begin{aligned}\nabla_{\mathbf{w}} f_{\text{MSE}}(\mathbf{y}, \hat{\mathbf{y}}; \mathbf{w}) &= \nabla_{\mathbf{w}} \left[\frac{1}{2n} \sum_{i=1}^n \left(\mathbf{x}^{(i)\top} \mathbf{w} - y^{(i)} \right)^2 \right] \\ &= \frac{1}{2n} \sum_{i=1}^n \nabla_{\mathbf{w}} \left[\left(\mathbf{x}^{(i)\top} \mathbf{w} - y^{(i)} \right)^2 \right] \\ &= \frac{1}{n} \sum_{i=1}^n \mathbf{x}^{(i)} \left(\mathbf{x}^{(i)\top} \mathbf{w} - y^{(i)} \right) \\ &= \frac{1}{n} \mathbf{X} (\mathbf{X}^\top \mathbf{w} - \mathbf{y})\end{aligned}$$

- Then plug in the current value of \mathbf{w} .
(Note that \mathbf{X} and \mathbf{y} are computed from the data and are constant.)

Gradient descent

- How *far* do we “move” left or right?
- Notice that, in the graph below, the **magnitude** of the slope (aka gradient) gives an indication of how far we need to go to reach the optimal **w**.



Gradient descent algorithm

- Set \mathbf{w} to random values; call this initial choice $\mathbf{w}^{(0)}$.

Python: `w = 0.01 * np.random.randn(M)` # Just an example!

Gradient descent algorithm

- Set \mathbf{w} to random values; call this initial choice $\mathbf{w}^{(0)}$.
- Compute the gradient: $\nabla_{\mathbf{w}} f(\mathbf{w}^{(0)})$

Gradient descent algorithm

- Set \mathbf{w} to random values; call this initial choice $\mathbf{w}^{(0)}$.
- Compute the gradient: $\nabla_{\mathbf{w}} f(\mathbf{w}^{(0)})$
- Update \mathbf{w} by moving opposite the gradient, multiplied by a **step size** ϵ .
$$\mathbf{w}^{(1)} \leftarrow \mathbf{w}^{(0)} - \epsilon \nabla_{\mathbf{w}} f(\mathbf{w}^{(0)})$$

Gradient descent algorithm

- Set \mathbf{w} to random values; call this initial choice $\mathbf{w}^{(0)}$.
- Compute the gradient: $\nabla_{\mathbf{w}} f(\mathbf{w}^{(0)})$
- Update \mathbf{w} by moving opposite the gradient, multiplied by a **step size** ϵ .
$$\mathbf{w}^{(1)} \leftarrow \mathbf{w}^{(0)} - \epsilon \nabla_{\mathbf{w}} f(\mathbf{w}^{(0)})$$
- Repeat...
$$\mathbf{w}^{(2)} \leftarrow \mathbf{w}^{(1)} - \epsilon \nabla_{\mathbf{w}} f(\mathbf{w}^{(1)})$$
$$\mathbf{w}^{(3)} \leftarrow \mathbf{w}^{(2)} - \epsilon \nabla_{\mathbf{w}} f(\mathbf{w}^{(2)})$$
$$\dots$$
$$\mathbf{w}^{(t)} \leftarrow \mathbf{w}^{(t-1)} - \epsilon \nabla_{\mathbf{w}} f(\mathbf{w}^{(t-1)})$$

Python: `w = w - EPS * gradient(w, X, y)`

Gradient descent algorithm

- How many iterations to run?
- Two alternative strategies:
 - Train for a fixed number of iterations T .

Gradient descent algorithm

- How many iterations to run?
- Two alternative strategies:
 - Train for a fixed number of iterations T .
 - Train until the difference in training cost diminishes below a threshold δ :

$$|f(\mathbf{w}^{(t-1)}) - f(\mathbf{w}^{(t)})| < \delta$$

Gradient descent demos

- 1-d
- 2-d