


CS 4342: Class 13

Jacob Whitehill

SVMs

Exercise

- Specify a hyperplane (\mathbf{w} , b) that would be used by an SVM for the following data:



Quadratic programming

SVM optimization problem

- Once again, we wish to:

- Minimize: $\frac{1}{2} \mathbf{w}^\top \mathbf{w}$

- Subject to: $y^{(i)} (\mathbf{x}^{(i)\top} \mathbf{w} + b) \geq 1 \quad \forall i$

- This is a **quadratic programming** problem: quadratic objective with linear inequality (and/or equality) constraints.
- There are many efficient solvers for quadratic programs.

Quadratic programming

- Quadratic programming is *not* a kind of computer programming.
- **Quadratic programming (QP)** problems are a kind of mathematical optimization problem:
 - Quadratic objective function (which we want to minimize or maximize).
 - Linear equality and/or inequality constraints.
- Same vein as linear programming, dynamic programming.

Quadratic programming

- Nonetheless, quadratic programs are typically *solved* using computer programs.
- As part of homework 4, you will use an off-the-shelf Python-based quadratic programming solver (**cvxopt**) to train an SVM.

cvxopt

```
cvxopt.solvers.qp(P,q[,G,h[,A,b[,solver[,initvals]]])
```

Solves the pair of primal and dual convex quadratic programs

$$\begin{array}{ll}\text{minimize} & (1/2)x^T P x + q^T x \\ \text{subject to} & Gx \preceq h \\ & Ax = b\end{array}$$

cvxopt

Quadratic objective

Linear objective

```
cvxopt.solvers.qp(P, q[, G, h[, A, b[, solver[, initvals]]])
```

Solves the pair of primal and dual convex quadratic programs

$$\begin{aligned} &\text{minimize} && (1/2)x^T P x + q^T x \\ &\text{subject to} && Gx \preceq h \\ & && Ax = b \end{aligned}$$

Inequality constraints

Equality constraints

- To train an SVM using a QP, we need to define the appropriate matrices from our training data.
- The \mathbf{x} comprises both the \mathbf{w} (hyperplane) and b (bias) (similar to how we implemented bias in linear regression).

cvxopt

Quadratic objective

Linear objective

```
cvxopt.solvers.qp(P, q [, G, h [, A, b [, solver [, initvals] ] ] ] )
```

Solves the pair of primal and dual convex quadratic programs

$$\begin{array}{ll} \text{minimize} & (1/2)x^T P x + q^T x \\ \text{subject to} & Gx \preceq h \\ & Ax = b \end{array}$$

Inequality constraints

Equality constraints

- **q** will just be 0 (we have no linear objective function).
- **P**: part of homework 4.

cvxopt

Quadratic objective

Linear objective

```
cvxopt.solvers.qp(P, q[, G, h[, A, b[, solver[, initvals]]]])
```

Solves the pair of primal and dual convex quadratic programs

$$\begin{array}{ll} \text{minimize} & (1/2)x^T P x + q^T x \\ \text{subject to} & Gx \preceq h \\ & Ax = b \end{array}$$

Inequality constraints

Equality constraints

- **G** and **h** need to encode the linear inequality constraints.
- We will not use **A** or **b** (optional parameters) since we have no equality constraints.

Defining G and h

- Suppose you have just two optimization variables — x_1 and x_2 — as well as the following constraints:
 - $2x_1 - 3x_2 \leq 2$
 - $x_1 + x_2 \geq 0$
- We need to express these both as linear inequality constraints (≤ 0).

Defining G and h

- Suppose you have just two optimization variables — x_1 and x_2 — as well as the following constraints:
 - $2x_1 - 3x_2 \leq 2$
 - $x_1 + x_2 \geq 0 \Rightarrow -x_1 - x_2 \leq 0$
- We need to express these both as linear inequality constraints (≤ 0).

Defining G and h

- Suppose you have just two optimization variables — x_1 and x_2 — as well as the following constraints:
 - $2x_1 - 3x_2 \leq 2$
 - $x_1 + x_2 \geq 0 \Rightarrow -x_1 - x_2 \leq 0$
- We need to express these both as linear inequality constraints (≤ 0).

$$\underbrace{\begin{bmatrix} 2 & -3 \\ -1 & -1 \end{bmatrix}}_{\mathbf{G}} \underbrace{\begin{bmatrix} x_1 \\ x_2 \end{bmatrix}}_{\mathbf{x}} \leq \underbrace{\begin{bmatrix} 2 \\ 0 \end{bmatrix}}_{\mathbf{h}}$$

More on accuracy metrics

True positives, false positives; true negatives, false negatives

- After training a classifier, we can apply the machine to new data to estimate the class label as \hat{y} .
- Examples:

- Step-wise classification:
$$g^{(j)}(\mathbf{x}) = \mathbb{I}[\mathbf{x}_{r_1, c_1} > \mathbf{x}_{r_2, c_2}]$$
$$\hat{y} = g(\mathbf{x}) = \mathbb{I} \left[\left(\frac{1}{m} \sum_{j=1}^m g^{(j)}(\mathbf{x}) \right) > 0.5 \right]$$

- The machine always outputs either 1 or 0.

True positives, false positives; true negatives, false negatives

- After training a classifier, we can apply the machine to new data to estimate the class label as \hat{y} .
- Examples:
 - Logistic regression: $\hat{y} = \sigma(\mathbf{x}^\top \mathbf{w})$
 - Although \hat{y} is probabilistic (in $(0,1)$), we can apply a **threshold** (e.g., $\tau=0.5$) to convert to a hard label:

$$\hat{y} = \mathbb{I}[\sigma(\mathbf{x}^\top \mathbf{w}) > \tau]$$

True positives, false positives; true negatives, false negatives

- Once both $y, \hat{y} \in \{0, 1\}$, we can compute the number of:
 - **True positives (TP)**: The number of positive examples ($y^{(i)}=1$) estimated by the machine to be positive ($\hat{y}^{(i)}=1$).
 - **False positives (FP)**: The number of negative examples ($y^{(i)}=0$) estimated by the machine to be positive ($\hat{y}^{(i)}=1$).
 - **True negatives (TN)**: The number of negative examples ($y^{(i)}=0$) estimated by the machine to be negative ($\hat{y}^{(i)}=0$).
 - **False negatives (FN)**: The number of positive examples ($y^{(i)}=1$) estimated by the machine to be negative ($\hat{y}^{(i)}=0$).

True positives, false positives; true negatives, false negatives

- We can also compute the *rate* of TP, FP, TN, FN:
 - **True positive rate:** The fraction of positive examples ($y^{(i)}=1$) estimated by the machine to be positive ($\hat{y}^{(i)}=1$).
 - **False positive rate:** The fraction of negative examples ($y^{(i)}=0$) estimated by the machine to be positive ($\hat{y}^{(i)}=1$).
 - **True negative rate:** The fraction of negative examples ($y^{(i)}=0$) estimated by the machine to be negative ($\hat{y}^{(i)}=0$).
 - **False negative rate:** The fraction of positive examples ($y^{(i)}=1$) estimated by the machine to be negative ($\hat{y}^{(i)}=0$).

True positives, false positives; true negatives, false negatives

- We can also compute the *rate* of TP, FP, TN, FN:
 - **True positive rate:** The fraction of positive examples ($y^{(i)}=1$) estimated by the machine to be positive ($\hat{y}^{(i)}=1$).
 - **False positive rate:** The fraction of negative examples ($y^{(i)}=0$) estimated by the machine to be positive ($\hat{y}^{(i)}=1$).
 - **True negative rate:** 1 - FPR
 - **False negative rate:** 1 - TPR

Example

- Suppose $\mathbf{y} = [1, 1, 0, 0, 1]$ and $\hat{\mathbf{y}} = [0.7, 0.3, 0.2, 0.9, 0]$.

Example

- Suppose $\mathbf{y} = [1, 1, 0, 0, 1]$ and $\hat{\mathbf{y}} = [1, 0, 0, 1, 0]$.
Apply threshold (0.5) to obtain labels in $\{0, 1\}$.

Example

- Suppose $\mathbf{y} = [1, 1, 0, 0, 1]$ and $\hat{\mathbf{y}} = [1, 0, 0, 1, 0]$.
- Then:
 - TPR =
 - FPR =
 - TNR =
 - FNR =

Different $\tau \Rightarrow$ Different TPR, TNR


- If you choose a different threshold τ , you will obtain different binary labels.
- Suppose $\mathbf{y} = [1, 1, 0, 0, 1]$. If the machine's real-valued outputs are $[0.7, 0.3, 0.2, 0.9, 0]$, then:
 - $\tau = -1 \Rightarrow \hat{\mathbf{y}} = [1, 1, 1, 1, 1] \Rightarrow \text{TPR}=1, \text{FPR}=1.$
 - $\tau = 0.19 \Rightarrow \hat{\mathbf{y}} = [1, 1, 1, 1, 0] \Rightarrow \text{TPR}=2/3, \text{FPR}=1.$
 - $\tau = 0.5 \Rightarrow \hat{\mathbf{y}} = [1, 0, 0, 1, 0] \Rightarrow \text{TPR}=1/3, \text{FPR}=1/2.$
 - $\tau = 0.7 \Rightarrow \hat{\mathbf{y}} = [0, 0, 0, 1, 0] \Rightarrow \text{TPR}=0, \text{FPR}=1/2.$

Different $\tau \Rightarrow$ Different TPR, TNR

- Higher threshold $\tau \Rightarrow$ lower TPR, lower FPR.
- Suppose $\mathbf{y} = [1, 1, 0, 0, 1]$. If the machine's real-valued outputs are $[0.7, 0.3, 0.2, 0.9, 0]$, then:
 - $\tau = -1 \Rightarrow \hat{\mathbf{y}} = [1, 1, 1, 1, 1] \Rightarrow \text{TPR}=1, \text{FPR}=1.$
 - $\tau = 0.19 \Rightarrow \hat{\mathbf{y}} = [1, 1, 1, 1, 0] \Rightarrow \text{TPR}=2/3, \text{FPR}=1.$
 - $\tau = 0.5 \Rightarrow \hat{\mathbf{y}} = [1, 0, 0, 1, 0] \Rightarrow \text{TPR}=1/3, \text{FPR}=1/2.$
 - $\tau = 0.7 \Rightarrow \hat{\mathbf{y}} = [0, 0, 0, 1, 0] \Rightarrow \text{TPR}=0, \text{FPR}=1/2.$


Receiver operating characteristics (ROC) curve

- If you plot the TPR v. FPR for *all possible thresholds*, you obtain the ROC curve of the classifier w.r.t. ground-truth:
 - $\mathbf{y} = [0, 0, 0, 0, 1, 1, 1]$ and $\hat{\mathbf{y}} = [0.1, 0.5, 0.8, 0.7, 0.9, 0.7, 0.95]$.




Receiver operating characteristics (ROC) curve

- If you plot the TPR v. FPR for *all possible thresholds*, you obtain the ROC curve of the classifier w.r.t. ground-truth:
 - $\mathbf{y} = [0, 0, 0, 0, 1, 1, 1]$ and $\hat{\mathbf{y}} = [0.1, 0.5, 0.8, 0.7, 0.9, 0.7, 0.95]$.



Area Under the ROC Curve (AUC)

- We can compute an aggregate metric over all possible thresholds by **integrating** the ROC curve:
 - $\mathbf{y} = [0, 0, 0, 0, 1, 1, 1]$ and $\hat{\mathbf{y}} = [0.1, 0.5, 0.8, 0.7, 0.9, 0.7, 0.95]$.



Area Under the ROC Curve (AUC)

- The AUC is a **threshold-independent** metric of how well the classifier discriminates between the 2 classes.
- The AUC of a classifier that always guesses correctly: 1
- The AUC of a classifier that always guesses incorrectly: 0
- The AUC of a classifier that always guesses the same value: 0.5
- The (expected) AUC of a classifier that guesses randomly: 0.5

Area Under the ROC Curve (AUC)

- The (expected) AUC is **not affected** by the ratio of positive to negative classes.
- AUC expresses how well a classifier can discriminate between two classes.
- Note that a classifier can have excellent *discriminability* but still make many *mistakes* in classification.

Area Under the ROC Curve (AUC)

- The AUC is also equivalent to the following:
 - Let (i,j) represent a randomly chosen pair of examples such that $y_i=1$ and $y_j=0$.
 - The AUC is the probability that $\hat{y}_i > \hat{y}_j$, i.e., the probability that the machine's output can correctly distinguish which example in the pair is positive.